



SAS Software Tour with iLink Mortgage, Inc.

SAS On-the-Job Activity

Purpose

This activity is a hands-on guide to using SAS Studio, Visual Analytics, and Model Studio within SAS Viya for Learners, while presenting learning in a real-world job situation. Students will modify and prepare data, create descriptive visualizations, estimate and compare machine learning models, and build and enhance visual reports.

SAS Software

We will use SAS Viya for Learners 4.0 (VFL4) software in this activity. Users can also use VFL3.5, but there may be slight differences in the user experience.

Industry Alignment

This On-the-Job activity fits within the mortgage industry of the financial sector. We explore data containing information about homeowner loan activity – and seek to predict which customers are most likely to default on their loans.

Table of Contents

SAS Software Tour with iLink Mortgage, Inc.	1
<i>Purpose</i>	1
<i>SAS Software</i>	1
<i>Industry Alignment</i>	1
Activity Notes and Requirements	3
<i>Learning Objectives</i>	3
<i>Estimated Completion Time</i>	3
<i>Experience Level</i>	3
<i>Prerequisite Knowledge</i>	3
Software	3
Content Knowledge	3
<i>Additional Notes</i>	3
Required Setup	3
Task 1: Examining Alex’s Work in SAS Studio	4
<i>Learning Objectives</i>	4
<i>Estimated Time of Completion</i>	4
<i>Task: Introduction to SAS Studio</i>	4
Task 2: SAS Visual Analytics – A Non-Coders Paradise	12
<i>Learning Objectives</i>	12
<i>Estimated Time of Completion</i>	12
<i>Task: Introduction to SAS Visual Analytics</i>	12
Task 3: Going Deeper with SAS Model Studio	32
<i>Learning Objectives</i>	32
<i>Estimated Time of Completion</i>	32
<i>Task: Going Deeper with SAS Model Studio</i>	32
Appendix	48
<i>Appendix A: Access Software</i>	48
<i>Appendix B: Helpful Documentation</i>	48
<i>Appendix C: Recommended Learning</i>	49

Activity Notes and Requirements

Learning Objectives

This activity provides the opportunity to learn and practice skills such as:

- Creating visualizations for descriptive statistics
- Modifying data to be used across SAS Viya for Learners applications
- Building and comparing machine learning models

Estimated Completion Time

This activity contains three tasks. All together this activity will take students roughly 90 to 120 minutes to complete.

Experience Level

This activity is designed for students new to SAS Viya. Students are encouraged to continue their SAS journey by visiting one of the multiple learning paths in *Appendix C: Recommended Learning*.

Prerequisite Knowledge

Software

No prior experience is needed with SAS Viya for Learners.

Content Knowledge

Students should have basic experience/knowledge of the following concepts:

- Statistical graphs like bar charts, histograms, and pie charts
- Statistical concepts like frequency, percentages, and mean (averages)

Additional Notes

Required Setup

Students need to complete each task sequentially to complete the activity as written.

Task 1: Examining Alex's Work in SAS Studio

Learning Objectives

This task provides the opportunity to learn and practice skills such as:

- Accessing and navigating the SAS Viya for Learners platform
- Learning to use Snippets in SAS Studio
- Running SAS Studio programs and examining results for key insights

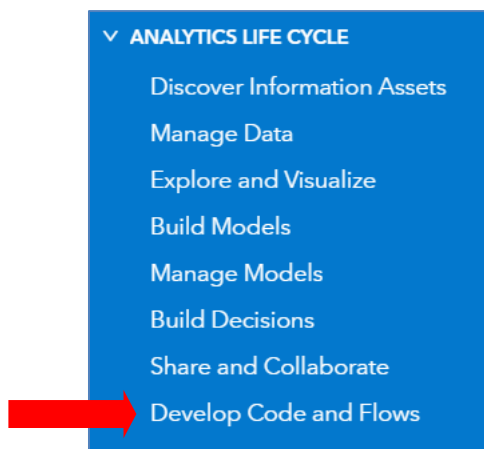
Estimated Time of Completion

This task is estimated to take about 30 minutes.

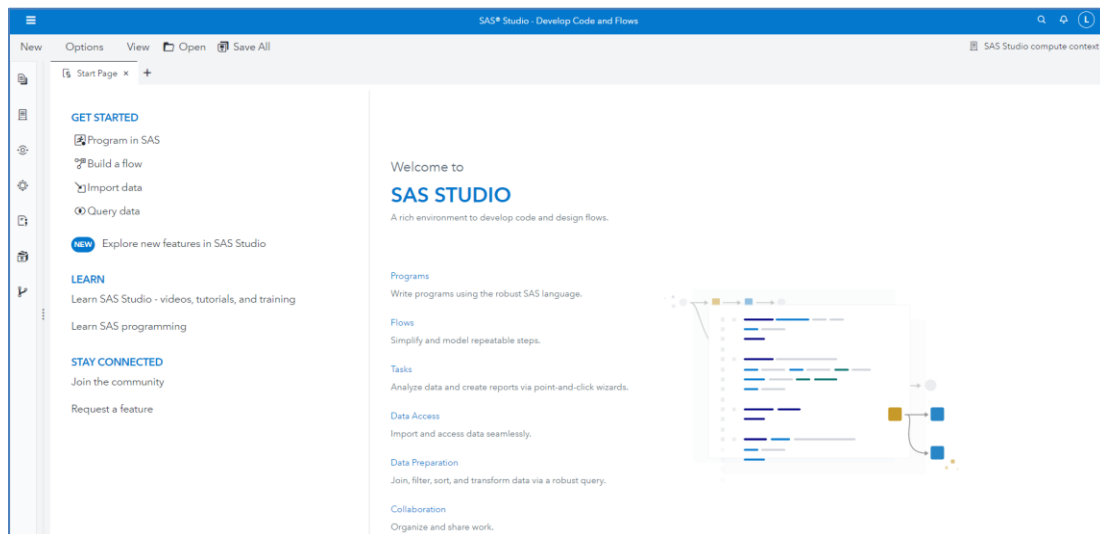
Task: Introduction to SAS Studio

Welcome to your first day at iLink Mortgage, Inc.! Your onboarding as a machine learning apprentice starts with an examination of your predecessor Alex's work. In particular, I'd like you to look at one of their SAS Studio projects, where they sought to predict which home loans were most likely to default. Defaulted loans cost our shareholders money, so we'll want any insights possible to determine which loans might be in trouble.

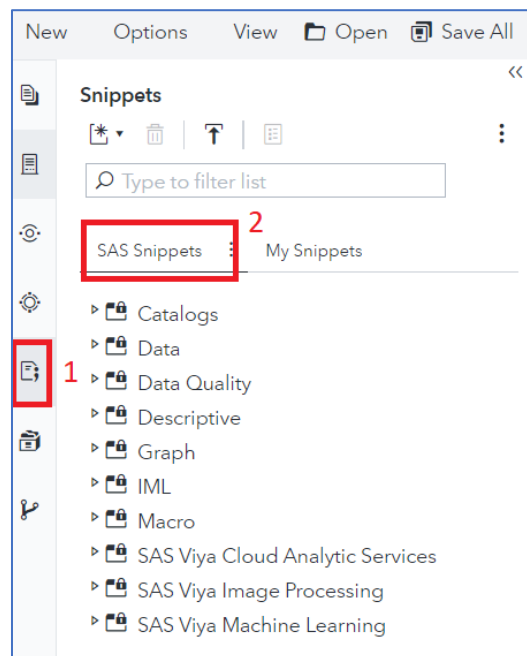
1. Log into SAS Viya for Learners.
2. You will be brought to SAS Drive, which is your starting point in SAS Viya for Learners.
3. From the top left corner, select the **Applications** menu ☰ and click **Develop SAS Code and Flows** from the Analytics Life Cycle options.



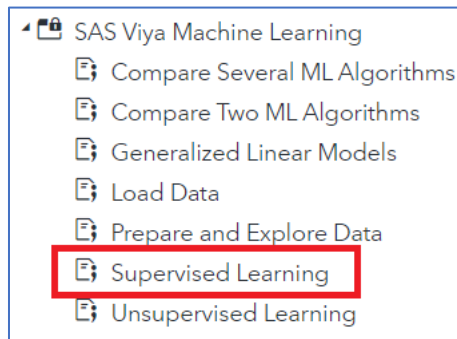
4. The **Welcome to SAS Studio** page should now be displayed.



5. From the **Start Page**, you can begin your analytical adventure in a variety of ways, from programming in SAS Studio to building flows. You can even embark on a learning journey with videos, tutorials, and training.
6. For this analysis, we'll simply access a program Alex saved as a **Snippet** (i.e., a chunk of reusable code). From the left pane, access the Snippets saved in your SAS Studio environment. SAS Snippets should be selected by default:



- Expand the section **SAS Viya Machine Learning** and open the **Supervised Learning** snippet:



8. The snippet starts with the following preamble:


```

1  /*****
2  /* This snippet showcases a sample Machine Learning workflow for          */
3  /* supervised learning using SAMPLEML.HMEQ data set. The steps include:  */
4  /* */
5  /* (1) PREPARE AND EXPLORE                                                */
6  /*   a) Load data set into CAS                                          */
7  /*   b) Explore                                                          */
8  /*   c) Partition                                                        */
9  /*   d) Impute                                                           */
10 /*   e) Identify variables that explain variance                        */
11 /* */
12 /* (2) PERFORM SUPERVISED LEARNING                                       */
13 /*   a) Fit model using random forest                                    */
14 /* */
15 /* (3) EVALUATE AND IMPLEMENT                                             */
16 /*   a) Score the data                                                  */
17 /*   b) Assess model performance                                         */
18 /*   c) Generate ROC and Lift charts                                    */
19 /*****/
20
21
22 /*****/
23 /* Define the macro variables for later use in the program              */
24 /*****/
25 /* Specify a folder path to write the temporary output files */
26 %let outdir = &_SASWORKINGDIR;
27
28 /* Create a CAS engine libref to save the output data sets */
29 %let caslibname = mycas;
30 libname &caslibname cas caslib=casuser;

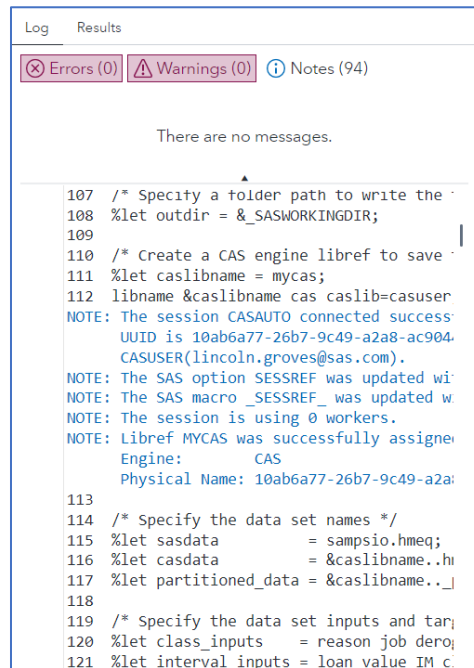
```

9. Please take a moment to read Alex's notes. This program does several things, as it takes raw data and uploads it into the SAS cloud, wrangles the data, then runs and analyzes a random forest model. The general analytical goal is to predict which loans will default, which occurs when the variable **BAD** = 1. BAD, in this case, is literally not good.

10. Once better understanding our analytical tasks, let the program rip! In the program window for *Supervised Learning.sas*, click on the **Run code** icon:

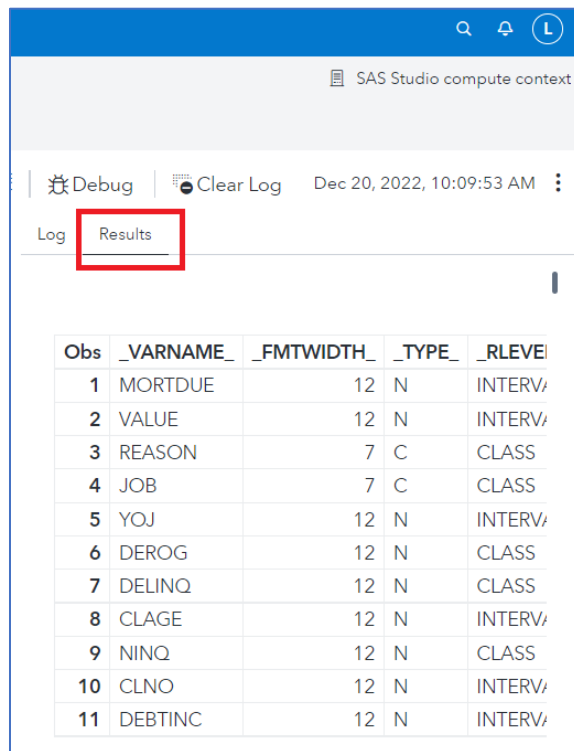


11. Examine the **Log** tab. Ensure that the program ran without any issues*:



* Check that the log contains no Errors or Warnings.

12. Now let's examine the results. From the rightmost pane, select **Results**:



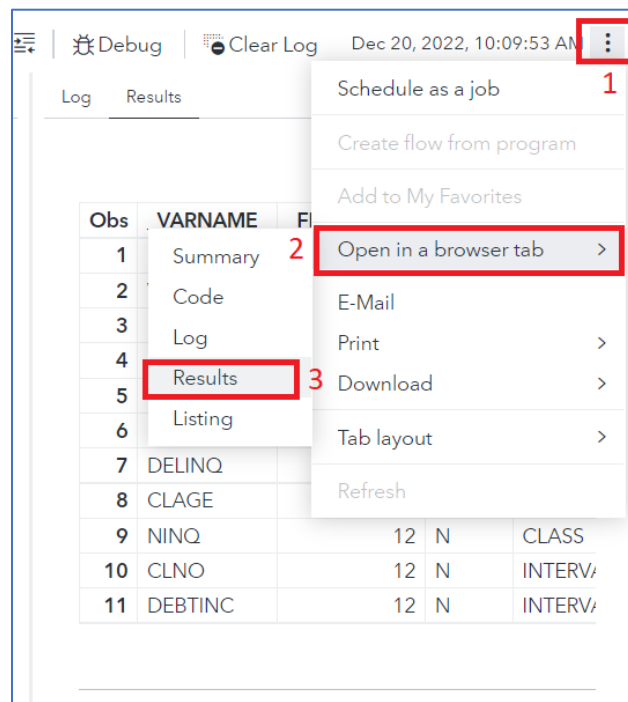
SAS Studio compute context

Debug Clear Log Dec 20, 2022, 10:09:53 AM

Log Results

Obs	_VARNAME_	_FMTWIDTH_	_TYPE_	_RLEVEL_
1	MORTDUE	12	N	INTERV
2	VALUE	12	N	INTERV
3	REASON	7	C	CLASS
4	JOB	7	C	CLASS
5	YOJ	12	N	INTERV
6	DEROG	12	N	CLASS
7	DELINQ	12	N	CLASS
8	CLAGE	12	N	INTERV
9	NINQ	12	N	CLASS
10	CLNO	12	N	INTERV
11	DEBTINC	12	N	INTERV

13. Those results feel a bit cramped, right? Let's open the **Results** window in a new browser tab. Follow the clicks below:



Debug Clear Log Dec 20, 2022, 10:09:53 AM

Log Results

Schedule as a job
Create flow from program
Add to My Favorites
Open in a browser tab
E-Mail
Print
Download
Tab layout
Refresh

Obs	VARNAME	F
1	Summary	
2	Code	
3	Log	
4	Results	
5	Listing	
6		
7	DELINQ	
8	CLAGE	
9	NINQ	12 N CLASS
10	CLNO	12 N INTERV
11	DEBTINC	12 N INTERV

14. Examine the model output, including how the data are transformed and which plots are used to assess the model's performance. In other words, scroll, and scroll some more!
15. Find and record the model **Fit Statistics**:

Fit Statistics					
Number of Observations	Squared Error			Mean Consequential Error	Multiclass Log Loss
	Divisor of Average	Average	Root Average		
4172	4172	0.062157	0.249313	0.081016	0.216747


Note: your results may differ marginally, as we didn't set a seed to ensure that the results were perfectly replicable. No worries though! This also illustrates that different training/validation/testing samples will produce slightly different results.

16. Now let's tweak the model slightly and see if we get a better Random Forest model. Return to your SAS program and find the **proc forest** code that starts on **line 128**. Locate the portion of the code that references the minimum leaf size for the individual trees within the forest model:

```

125  /****** */
126  /* Build a predictive model using Random Forest */
127  /****** */
128  proc forest data=&caslibname.._prepped ntrees=50 numbin=20 minleafsize=5;
129      input &interval_inputs. / level = interval;
130      input &class_inputs. / level = nominal;
131      target &target / level = nominal;
132      partition rolevar=_partind_(train='1' validate='0');
133      code file="&outdir./forest.sas";
134      ods output FitStatistics=fitstats;
135  run;

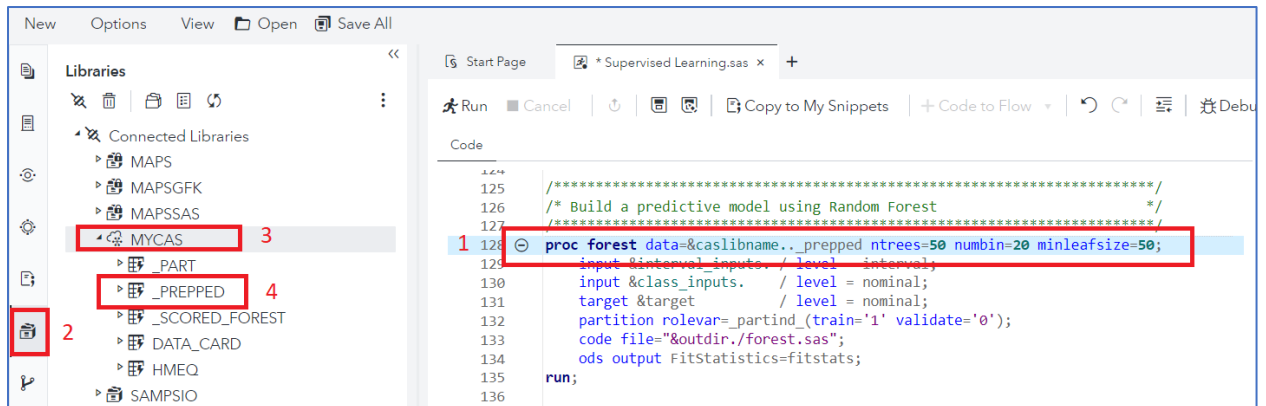
```

17. Let's change **minleafsize=50**. Then resubmit the code using  **Run**.
18. Does this model fit the data better, as indicated by the average squared error? My findings:

Fit Statistics					
Number of Observations	Squared Error			Mean Consequential Error	Multiclass Log Loss
	Divisor of Average	Average	Root Average		
4172	4172	0.097346	0.312004	0.128236	0.316123

19. Turns out that increasing the minimum leaf size to 50 does not improve model fit in this scenario. Can't win them all.

20. Finally, we would like to use the data created from this program in two other SAS applications: **SAS Visual Analytics** and **SAS Model Studio**. Since we don't want to spend much time cleaning the data again, let's examine which datasets have been saved for us, by first looking at the dataset used in the forest modeling.
21. While I don't expect you to read SAS code already – particularly those crazy macros and the "&" notation – we can see from the prompts in the screenshot below that the table name is **_prepped**. Moreover, the data reside in the SAS Library **MYCAS**. The evidence, by numbers:



Note: Step 1 is in the programming window. Steps 2-4 are focused on the Libraries pane.

22. We'll need to save and promote **_PREPPED** to the shared CAS folder, so that it can be accessible in SAS Visual Analytics. Again, the goal here isn't to turn you into a full-fledged hacker on your first day, so we've written the following code for you to use:

```

/*****
/* Prepare Data for use in SAS VA
/*****

/* To save the table to a physical copy*/
proc casutil;
    save casdata='_prepped' incaslib='casuser'
    casout='_prepped.sashdat' outcaslib='casuser' replace;
quit;

```

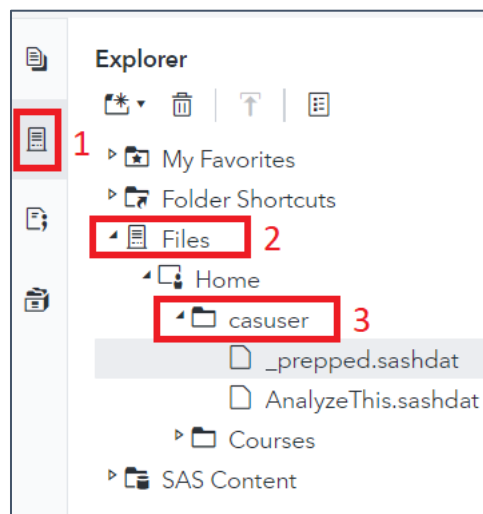
23. Add the code above to the end of your program:

```
208 /*****  
209 /* Prepare Data for use in SAS VA */  
210 /*****  
211 /* To save the table to a physical copy*/  
212 ⊖ proc casutil;  
213     save casdata='_prepped' incaslib='casuser'  
214     casout='_prepped.sashdat' outcaslib='casuser' replace;  
215 quit;  
216
```

24. Without any subsection of the code selected, resubmit the entire program via  **Run**.

And, voila, **_prepped** should now be accessible in SAS Visual Analytics!

25. We'll do a quick check to ensure that **_prepped** is in your CASUSER folder. Follow these clicks to investigate:



26. Congratulations on finishing the first part of our mortgage adventure!

Task 2: SAS Visual Analytics – A Non-Coders Paradise

Learning Objectives

This task provides the opportunity to learn and practice skills such as:

- Navigating Visual Analytics in the SAS Viya for Learners platform
- Learning to use the Auto Chart feature
- Creating visualizations using different types of variables (categorical, numeric, geographic)
- Using predictive modeling functions within the Visual Analytics application

Estimated Time of Completion

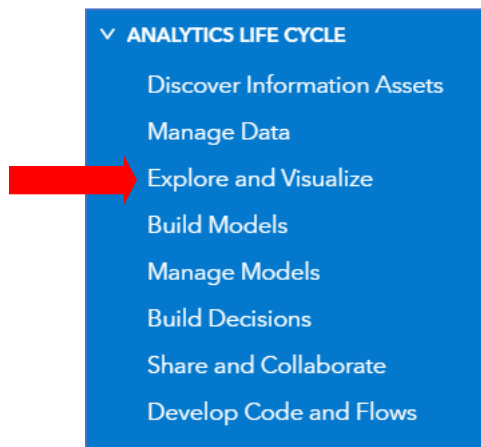
This task is estimated to take about 30 minutes.

Task: Introduction to SAS Visual Analytics

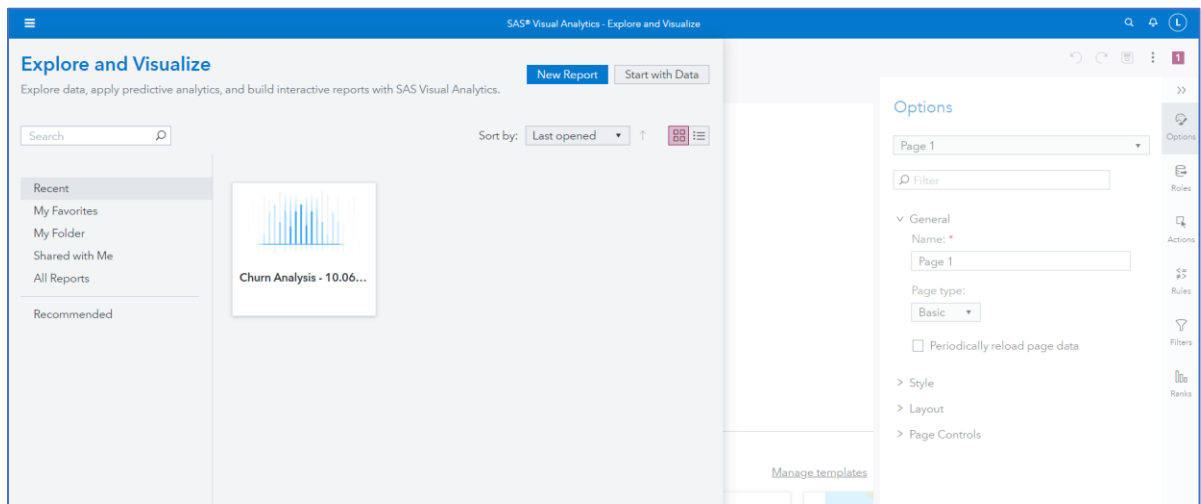
Coding is great. But as a new employee of iLink Mortgage, Inc. – and new user of SAS – you’d likely prefer to start your analytical journey with some no-code options. Fortunately, SAS Visual Analytics has several helpful tools to guide you along that path!

For this task, you’ll create your first predictive models to examine whether the loan will default (i.e., BAD = 1). Buckle up for the adventure!

1. Let’s move over into SAS Visual Analytics. Return to the **Applications** menu (≡) in the upper left corner and select **Explore and Visualize** under the *Analytics Life Cycle* options.



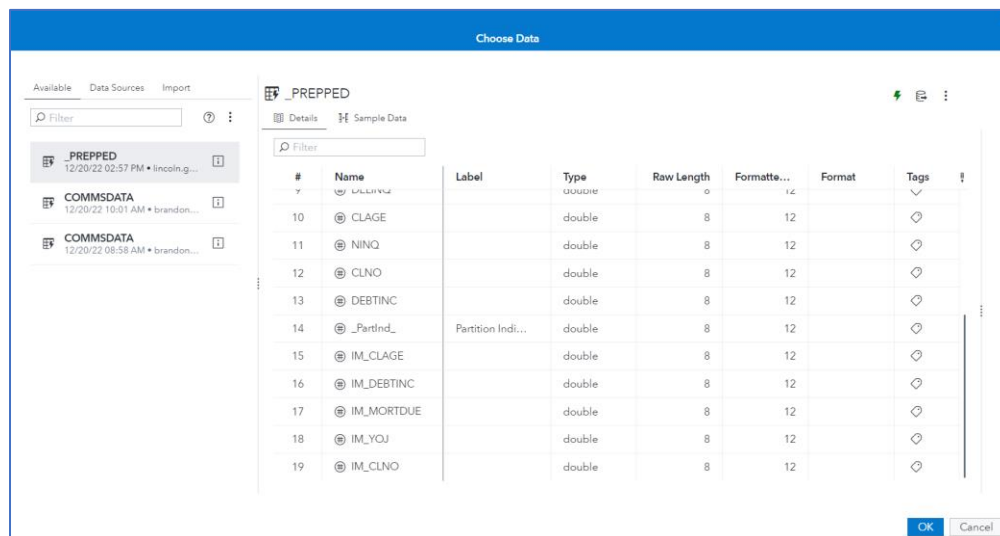
2. Welcome to **SAS Visual Analytics**! While your view will differ marginally, it should appear roughly akin to the following:



- To continue our SAS VA adventure, choose **Start with Data** from the main **Explore and Visualize** window:

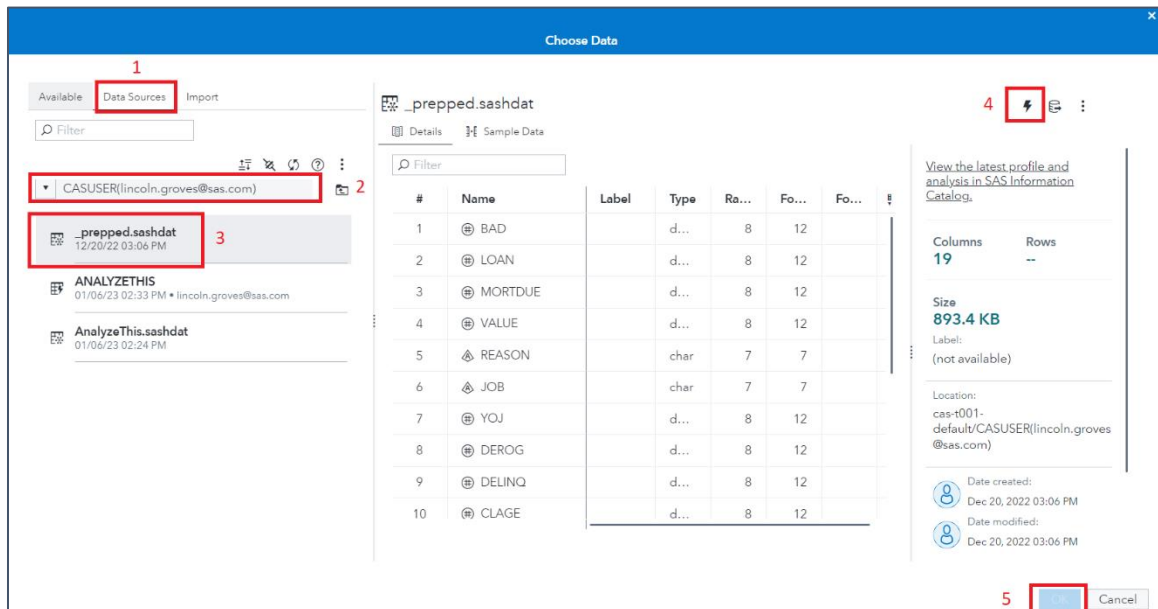
Start with Data

- And... this is the part where you try to remember, and find, the dataset used for modeling in SAS Studio. No worries if you forgot that name – it's **_PREPPED**. And it could be one of the available in-memory tables:



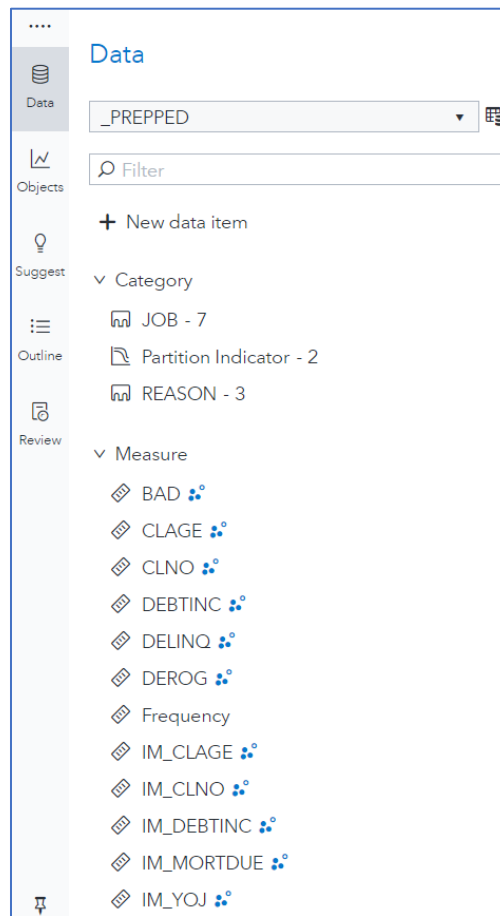
- However, if this dataset is not already loaded into memory, then select the **Data Sources** tab. Then select the default folder and drill down to your CASUSER folder. **_PREPPED.sashdat** is there – but not yet loaded into memory (which is why you didn't

see it above). How do we get data into memory? Click the **Load into Memory** button – i.e., the lightning bolt – then click **OK**. Follow these steps:

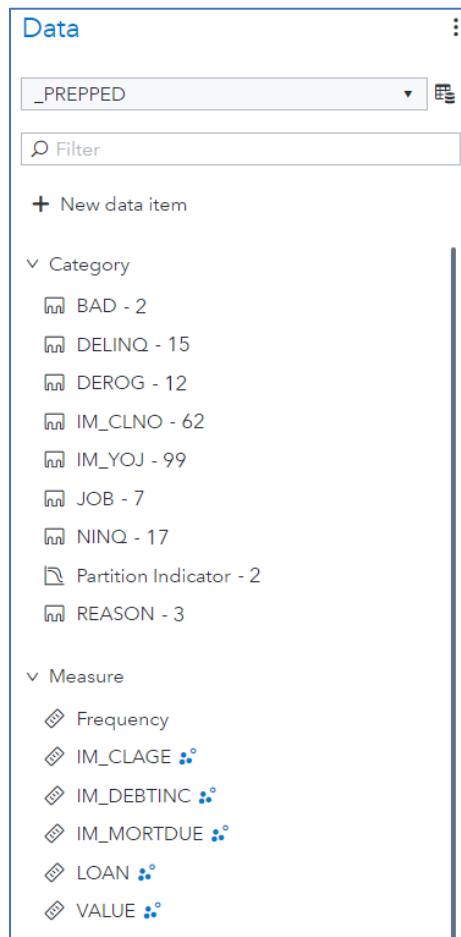


6. While a lot of the metadata definitions transferred from the SAS Studio environment, some of them did not. By cleaning up the metadata in SAS VA, we can explicitly state which variables are categorical variables, which are inputs on an interval, and which variables should be ignored all together. Let's get started...


Click the **Data** button on the left-hand side pane and examine the current variables in your analysis:



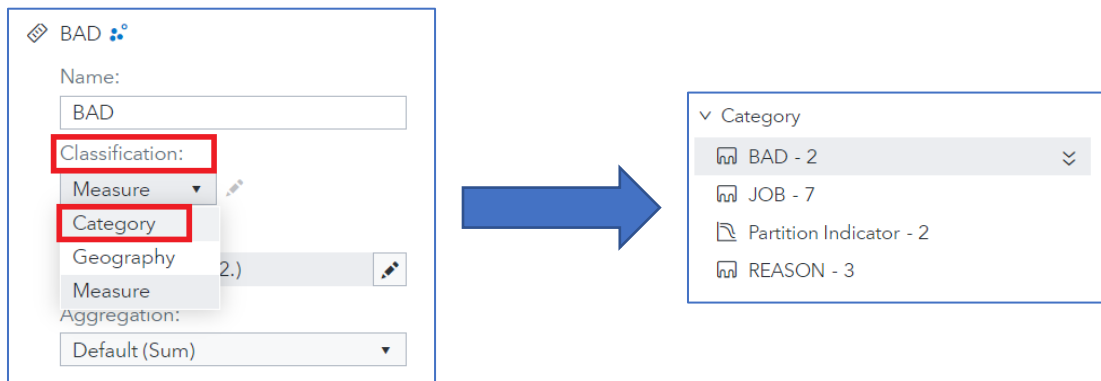
7. Whew! There are a lot of variables! For our upcoming modeling, we'll want to match the following metadata definitions:



8. And how do we do that, you might ask? Each variable has an **Edit properties** button which allows you to change the metadata for that variable:



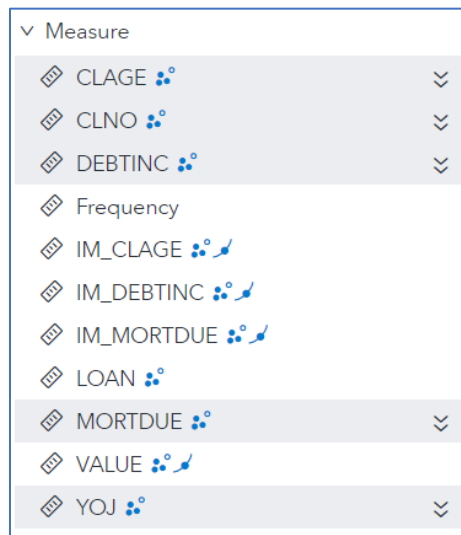
9. For our analysis, we want to ensure that BAD, DELINQ, DEROG, IM_CLNO, IM_YOJ, JOB, NINQ, and REASON are all classified as **Category**. Why? Because we'd like to treat them as Yes/No variables, instead of the usual 1/0 binary variable that is coded.
10. *BAD*, our outcome variable in the modeling, is possibly still classified as a **Measure**. If this is true, we can change this by clicking on the **Edit properties** button and then changing the classification from **Measure** to **Category**. The variable will now appear under **Category** and will show the number of unique levels in the data:



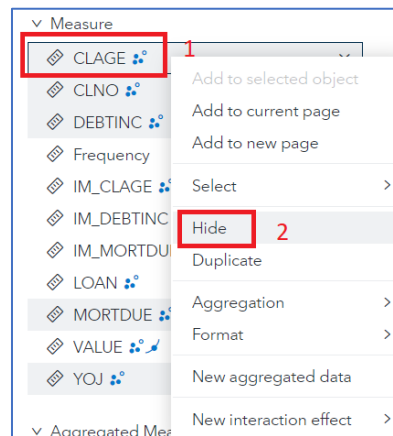
11. If required, make the change from **Measure** to **Category** for these other variables:
DEROG, DELINQ, NINQ, IM_CLNO, and IM_YOJ
12. Finally, we'd like to exclude some of the Measures – as these variables will not be used in our modeling. Before getting started, ensure that no variables are selected. Why? Because those variables will be included in the action. If they are, click on the “Clear Selection” option at the bottom of the data pane:




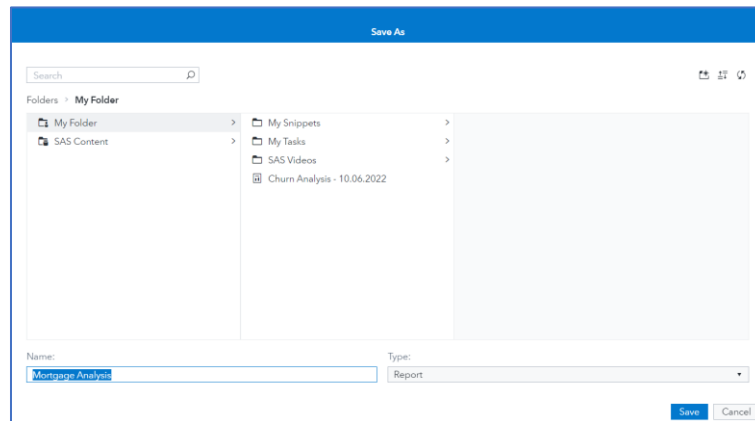
13. Then, select CLAGE, CLNO, DEBTINC, MORTDUE, and YOJ as shown:



14. Next, right-click on one of the selected variables, then select **Hide**:

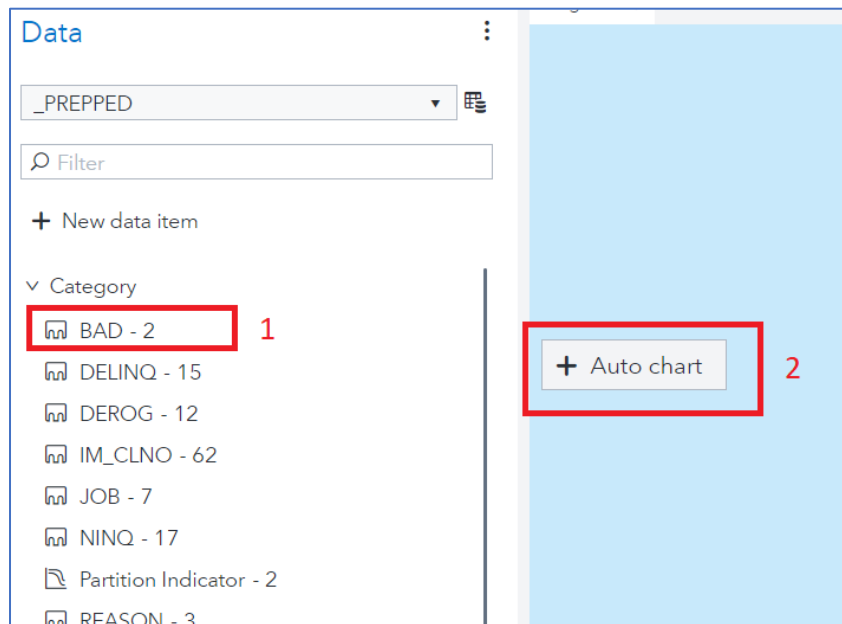


15. Presto! Variables be gone – and your metadata should now match the screenshot from above.
16. This is a great time to save your program. Find the classic **floppy-disk icon**  in the upper-right hand corner. Name the analysis something creative like **Mortgage Analysis** and **Save** it in the default location:




Please continually save your work as the analysis progresses. Lost work is kind of a bummer 😞!

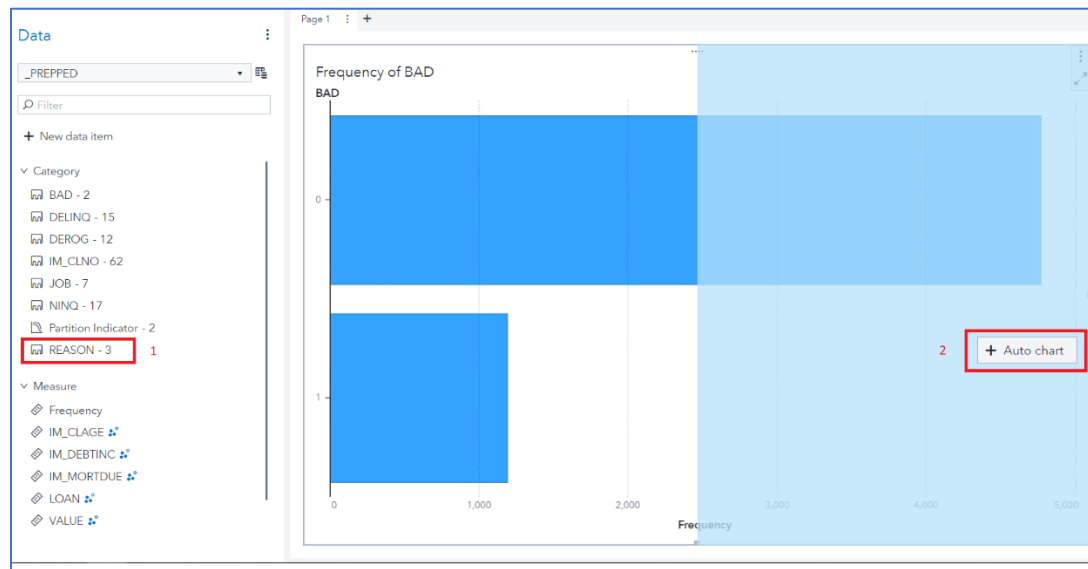
17. Now that the VA report is established and the data are where they need to be, let's better understand select variables by using **Auto charts**. With the **Data** icon selected, click on *BAD* and then drag and drop the variable to the canvas:




18. With **Auto chart**, Visual Analytics will automatically intuit what you'd like plotted. Here it will produce a bar chart with frequencies:



19. We'd also like to see distributions for five other variables: REASON, JOB, LOAN, IM_DEBTINC, and IM_MORTDUE. To add another variable to the canvas, move it from the **Data** tab to a unique space on the canvas. You'll see the **+Auto chart** option when the space is available for a new chart:




20. Variables can be placed above/below or left/right of existing objects. And if you make a mistake then no worries – there is an undo button in the upper-right hand corner:

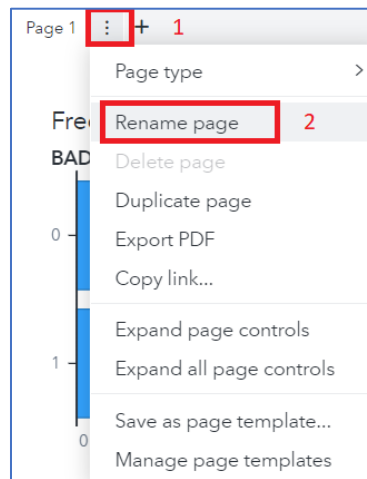


Use it. And use it often when needed.

21. Now replicate the following:



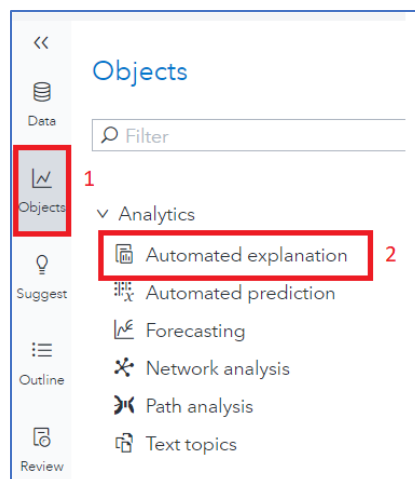
22. Change the name of “Page 1” to “DescriptiveStatistics”, by clicking on the **Page menu** and then selecting **Rename page** (or you can simply double click on the page tab):



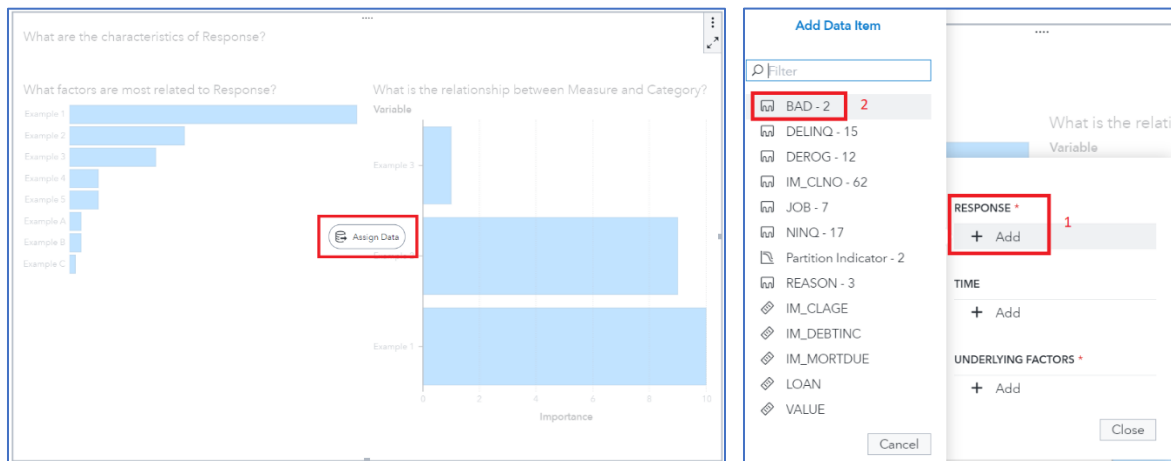
23. Create a **New page** by clicking the + button next to the “DescriptiveStatistics” tab. In this new canvas, rename “Page 2” to “AutomatedExplanation”.



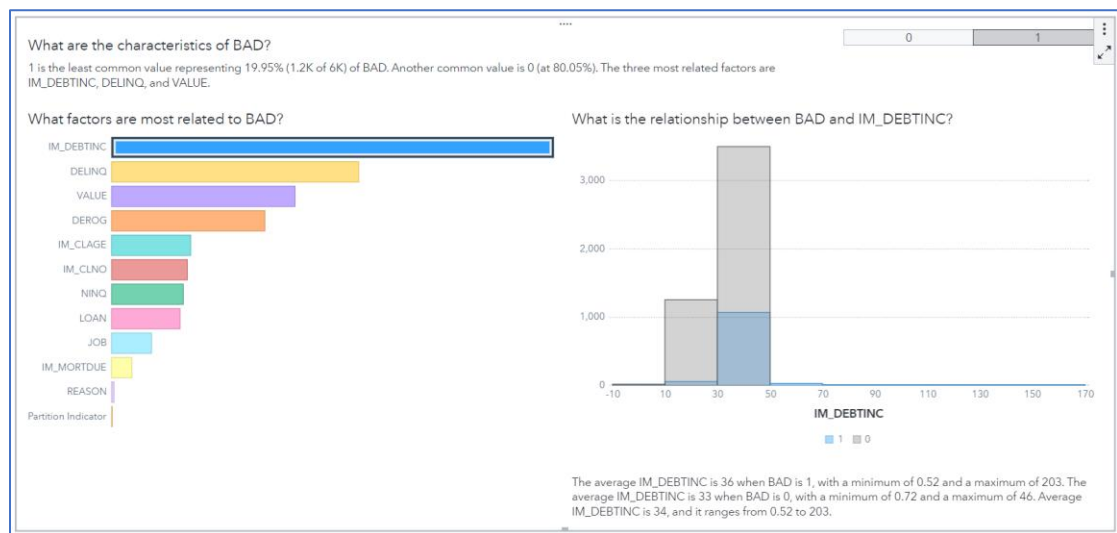
24. Let's add an **Automated Explanation** object to this new canvas. It can be found under **Objects** under the left-hand pane and then **Analytics**:



25. Next, click on **Assign Data** in the **Automated Explanation** object and then choose *BAD* as the **Response**:



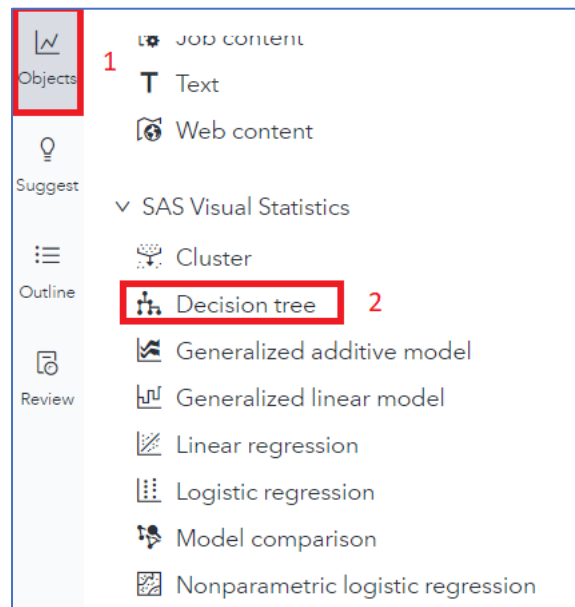
26. Examine the output from the **Automated Explanation** object. Start by clicking on **BAD=1** in the upper-right and then explore how individual factors are related to *BAD*, as follows:



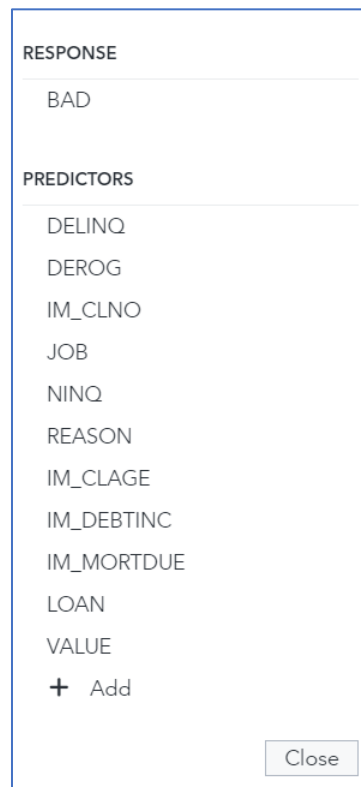
27. While **Automated Analysis** is a fast way to understand the bi-variate relationship between variables, we have more complex modeling tools. In other words, we have much cooler stuff. Let's start with a **Decision Tree** and then build more sophisticated models from there.

28. Create a **New page** and change the name to "DecisionTree".

29. From the **Objects** pane on the left, find the **Decision Tree** object under **SAS Visual Statistics** and drag-and-drop it on the DecisionTree canvas:

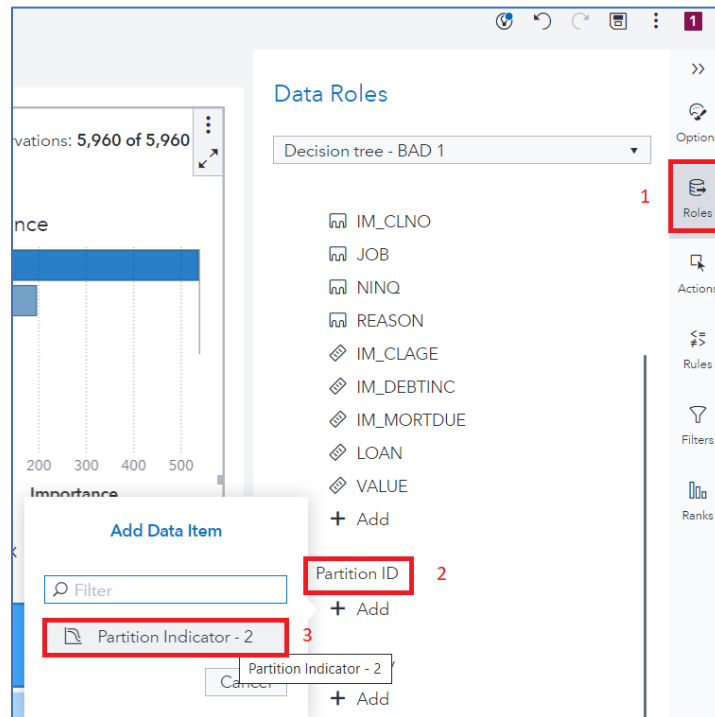


30. It's again time to **Assign Data**. Select **BAD** as the **Response** and the following variables as the **Predictors**. Then click **Close**:



31. Now – because we're aspiring data scientists, after all – let's add a data partition to allow for honest assessment. First, check to see if you have a partition already defined. From the

right-pane, select **Roles**. Then, scroll down to **Partition ID**. See if the **Partition Indicator** variable exists:



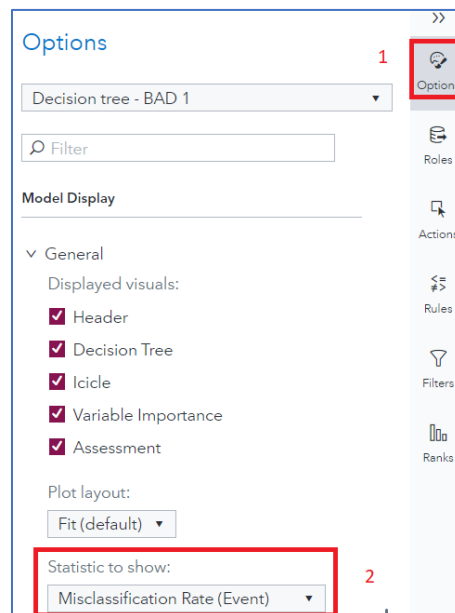
32. If yes, then choose *Partition Indicator* as your Partition ID – and skip the next two steps.
If not, you'll need to create a partition. No worries, we've got you covered.
33. For the unlucky few out there, Partition Indicator was not imported as an actual partition. To fix this – by creating a new variable – click on **Data** → **+ New data item** → **Partition**. Match the setting below in the pop-up window, and then click OK.

The 'New Partition' dialog box is shown. It has a blue header with the title 'New Partition'. Below the header, there are several fields and options:

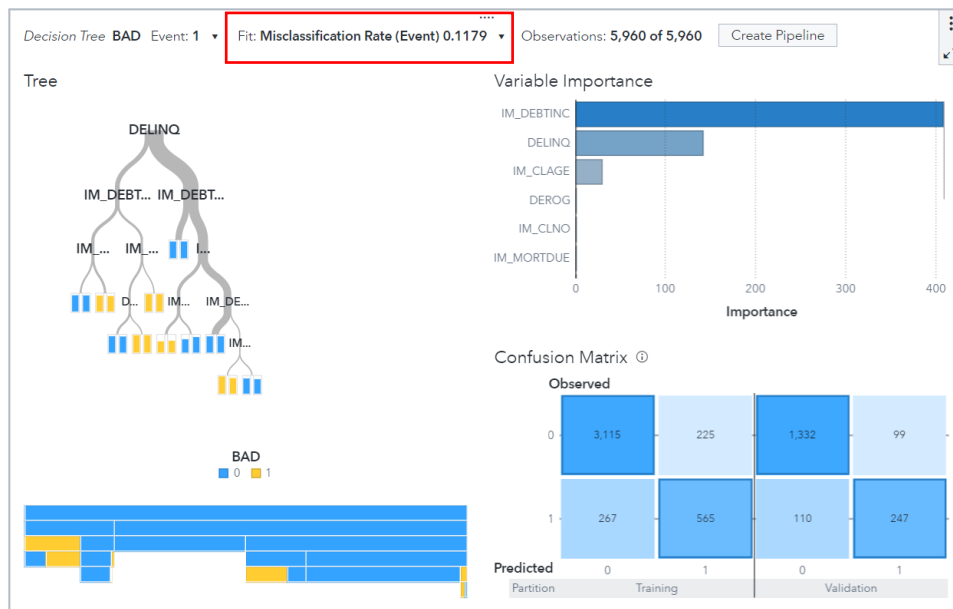
- Based on:** Two radio buttons are present: 'Data item' (selected) and 'Sampling'.
- Data item:** A dropdown menu showing 'Partition Indicator'.
- Number of partitions:** A dropdown menu showing '2'.
- Validation data value:** A dropdown menu showing '0'.
- Training data value:** A dropdown menu showing '1'.

At the bottom right, there are two buttons: 'OK' (highlighted in blue) and 'Cancel'.

34. You can then use the newly created partition variable as, yes, your partition. And you're all caught up.
35. One last item before we examine our output – let's explore how to use other model goodness-of-fit statistics. **Misclassification Rate** is the ratio of misclassified observations to total observations and is a common selection statistic used by many data scientists. Let's explore that option here – with the general underlying point that we have many statistics to choose from when determining the "optimal" model.
36. From the right-pane, select **Options** and then scroll down until you see the section for **Model Display**. Under **General**, change the **Statistic to show** to **Misclassification Rate (Event)** (**Event**):



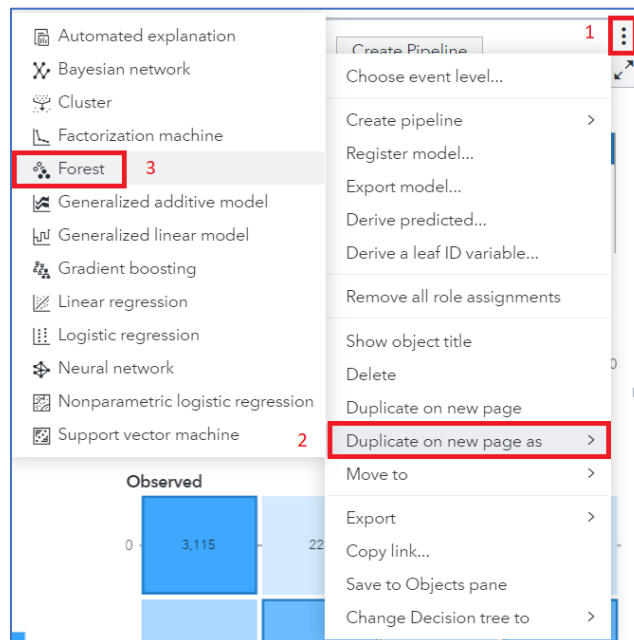
37. Observe and record the Misclassification Rate. Did you get a similar finding?



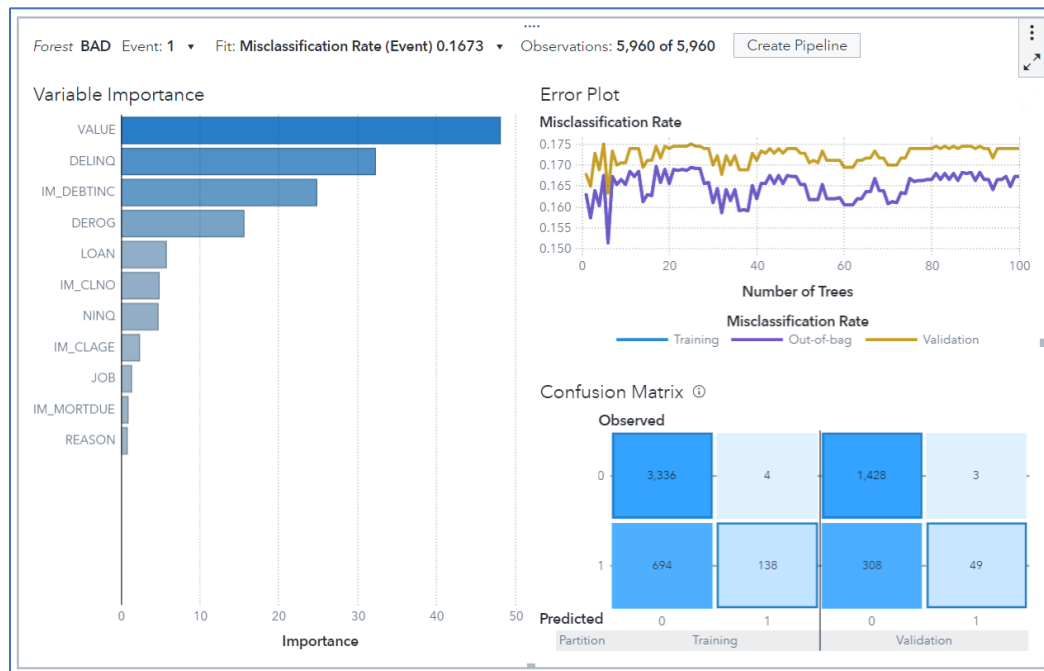
Don't sweat it if you didn't get the exact results reported above. It's likely because we're using different samples.

38. Now that was just a single model – based upon a single decision tree. What is often better than a single tree? A forest! Let's use the information from the Decision Tree to quickly run a couple of other models.

From the **Object menu for the Decision tree**, hold down **ALT** and click **Duplicate on new Page as > Forest**. Follow through with these clicks:



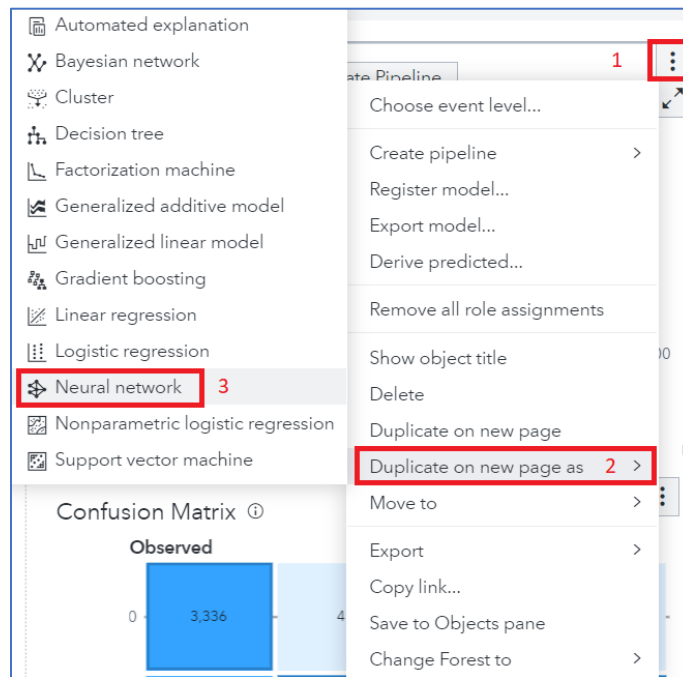
39. This great hack takes all the metadata definitions from the decision tree and puts them into the forest model. And the model is run instantaneously! As in the previous example, change the **Statistic to show** to **Misclassification Rate (Event)** and then examine the output:



40. Based upon **Misclassification Rate (Event)** is this model better or worse than the decision tree?

...Yup, it's worse. Weird. But that just means that we probably need to tune the hyperparameters – which is a bedtime story for another time.

41. Rename "Page 4" to "Forest" and let's run one more model – a Neural Network! From the **Object Menu for Forest**, again use **Alt + Duplicate on new page as > Neural Network** to create a third competitor model in a new page:

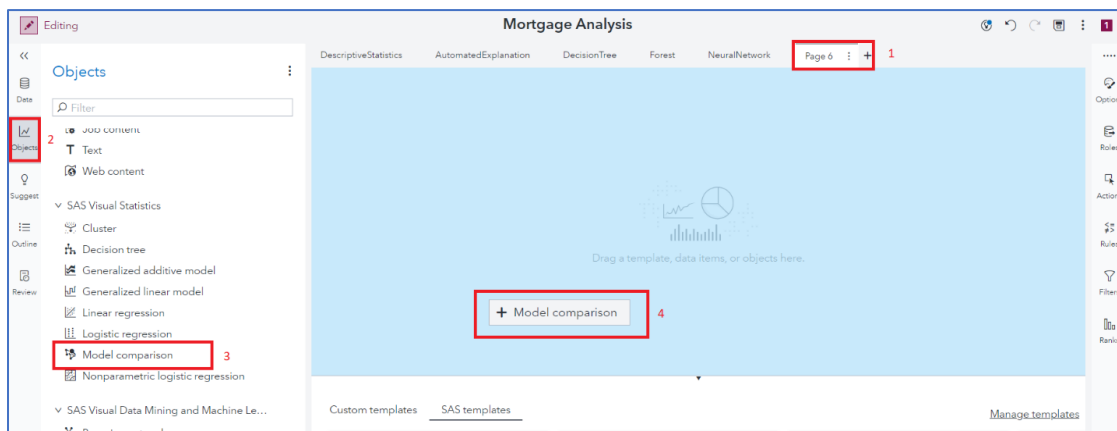


42. To make the models comparable, let's again change the **Statistic to show to Misclassification Rate (Event)**. Also, if you haven't already, let's change the name of the new page to "NeuralNetwork" and then examine the results:

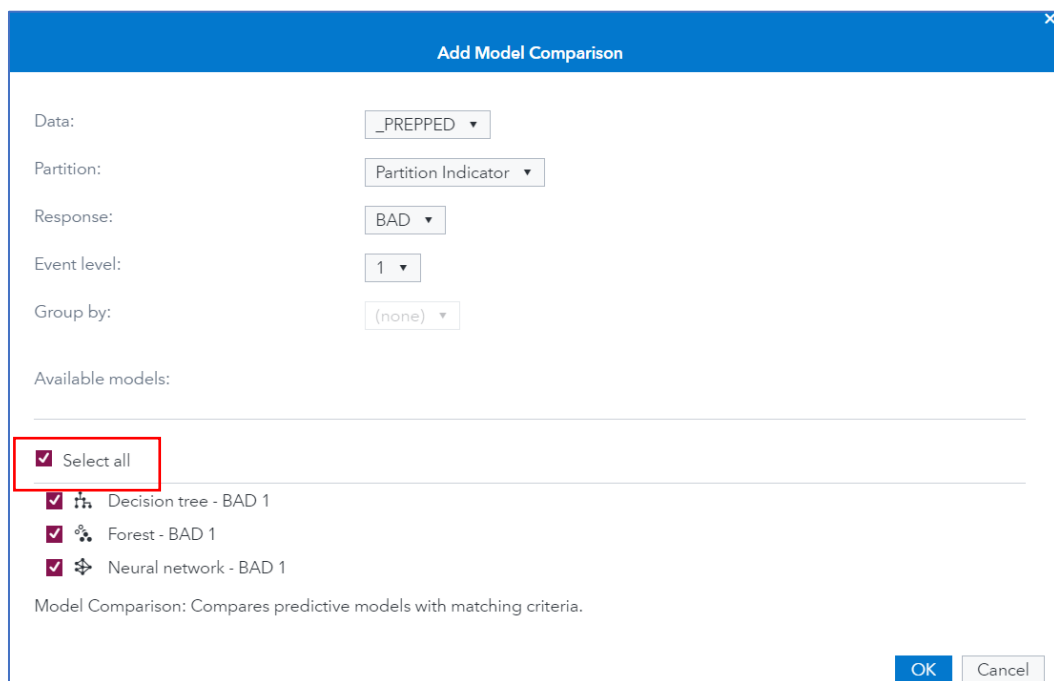


Note: it looks like we lost some observations – likely because neural networks use complete case analysis and we have missing data. But, this is a problem for another day.

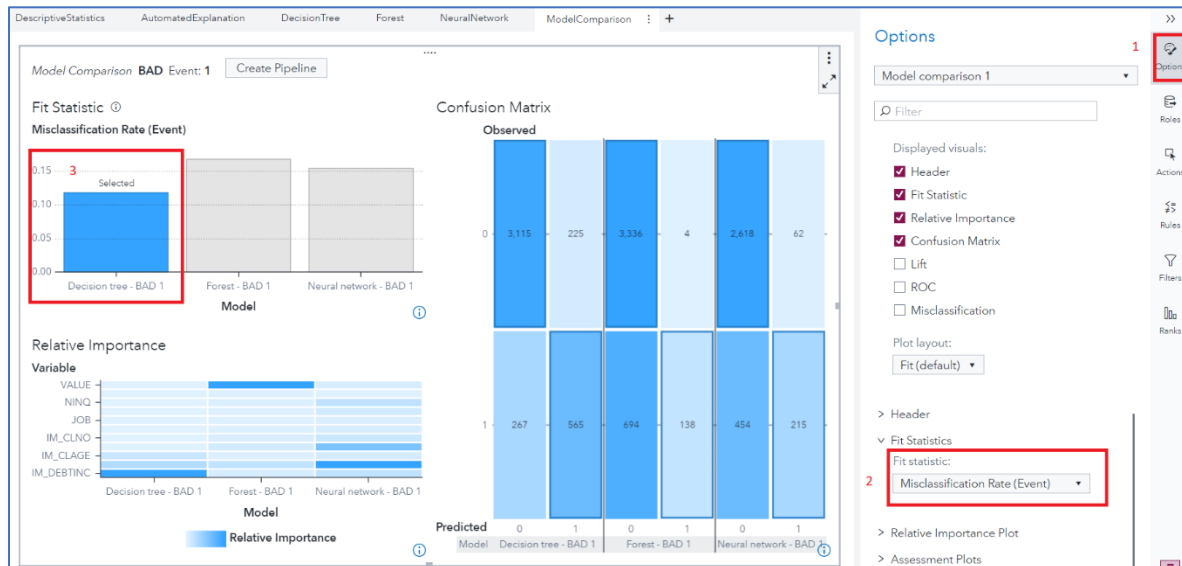
43. In general, how did the neural network do? Well, again, not great. The misclassification rate is 0.1541, which is – again – higher than the simpler decision tree model.
44. The Decision Tree model appears to have the lowest Misclassification Rate for the current set of models. So, we can crown it as our initial “champion” model. But, before moving on, let’s create one more page and use the **Model Comparison** object to nicely summarize the findings.
45. To do so, create a new page, find the **Model Comparison** object and drag it to the new canvas:




46. The **Add Model Comparison** window appears. Use the default settings, **Select all** models, and then click **OK**:




47. While data are being analyzed, rename “Page 6” to “ModelComparison”.
48. Finally – with the **Model Comparison** object selected – click **Options** and change the **Fit Statistic** to **Misclassification Rate (Event)**. Do you also get the Decision Tree as the champion?



49. Hopefully there are no surprises there – as the Model Comparison object simply replicates and combines the findings from our three predictive models.
50. Our last Visual Analytics task is to export the champion model to **SAS Model Studio** – where we have many more no-code options to explore. To do so, return to the DecisionTree page. Find the following:



51. The **Create Pipeline** button has been hiding in plain sight all along! Pipelines are a fundamental tool in SAS Model Studio, which you will learn a lot about in the next section. Click that **Create Pipeline** button and then select **Add to new project**:



52. Your data and metadata definitions will then be transferred to SAS Model Studio, where the next adventure awaits!

Task 3: Going Deeper with SAS Model Studio

Learning Objectives

This task provides the opportunity to learn and practice skills such as:

- Navigating the SAS Model Studio application
- Creating and running pipelines for supervised ML models
- Incorporating open-source tools
- Comparing and assessing models

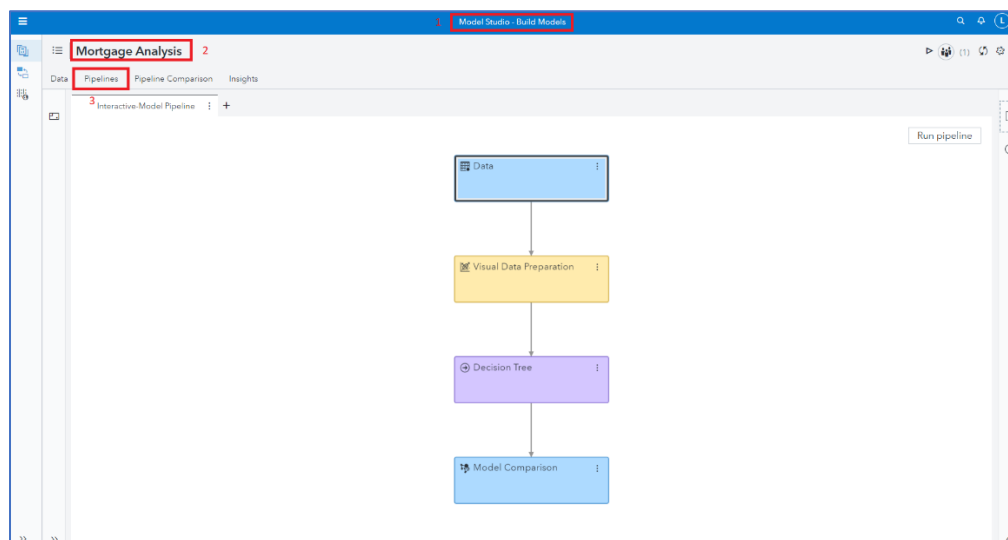
Estimated Time of Completion

This task is estimated to take about 45 minutes.

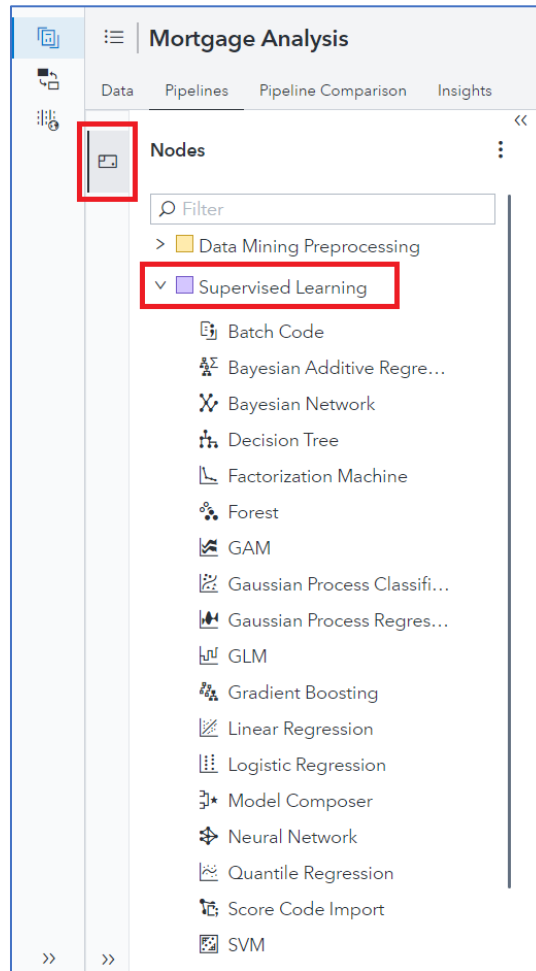
Task: Learn to love SAS Model Studio

SAS Visual Analytics is a great starting point for those new to SAS and analytics. But SAS VA simplifies many of the modeling options to make it a more user-friendly experience. In other words, there is a tradeoff between ease-of-use and modeling options available. With SAS Model Studio, you get a much more robust no-code and (limited) code environment – but it takes a bit more technical know-how. So, let's take our analytical endeavors to the next level with this great tool.

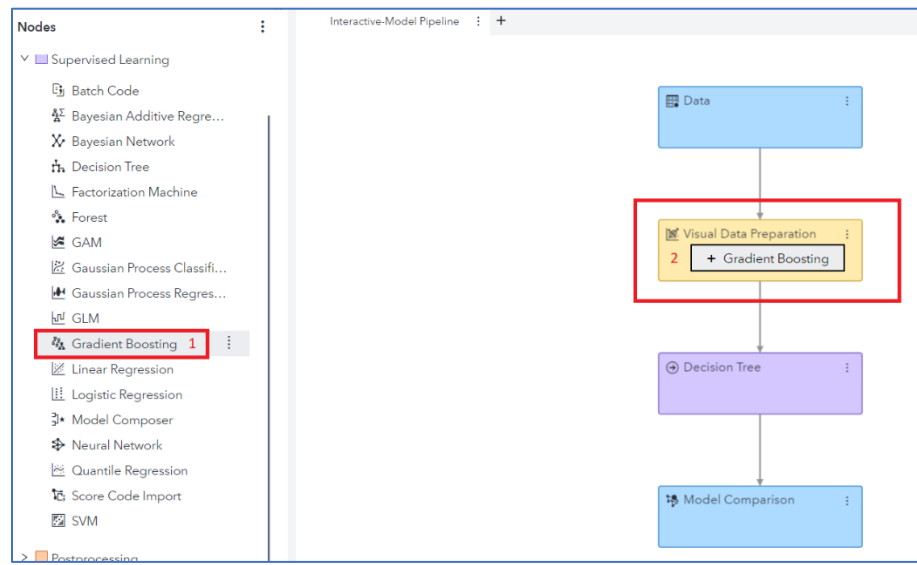
1. If you exported the pipeline properly in the last section, you should have a new SAS Model Studio Project named "Mortgage Analysis" – and the project should open to the **Interactive Model Pipeline**:



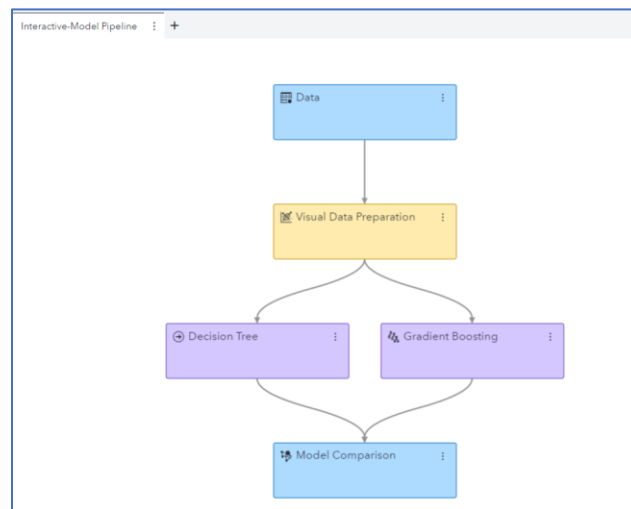
2. SAS Model Studio is pipeline centric and you can see our champion decision tree from our SAS Visual Analytics investigation has been converted to a pipeline.
3. Before running our first SAS Model Studio pipeline, let's explore how easy it is to add another model to the pipeline. Locate the **Nodes** button in the left-hand corner. Expand the **Supervised Learning** nodes and see what is available:




4. Now that's a lot of options! Let's see if the default **Gradient Boosting** model can defeat our Decision Tree. Grab the Gradient Boosting node and drop it on top of the **Visual Data Preparation** node:



5. The new pipeline appears as:



6. Submit the pipeline by selecting Run pipeline in the upper right-hand corner.
7. While this pipeline is running, I can show you one of my favorite features of SAS Model Studio. It's in the prebuilt pipelines, which makes running many models very easily. Create a new pipeline by clicking on the + next to the "Interactive-Model Pipeline" tab:



8. The **New Pipeline** window appears. Under **Select a pipeline template**, choose **Browse**:

New Pipeline

Name *

Pipeline 1

Description:

☒ Select a pipeline template
 1

Blank template

Browse

☐ Automatically generate the pipeline ⓘ
 2

☒ Set automation time limit ⓘ

10 minutes

Advanced Settings

Save

Cancel

9. Check out all the pipeline templates available!

Browse Templates

Filter

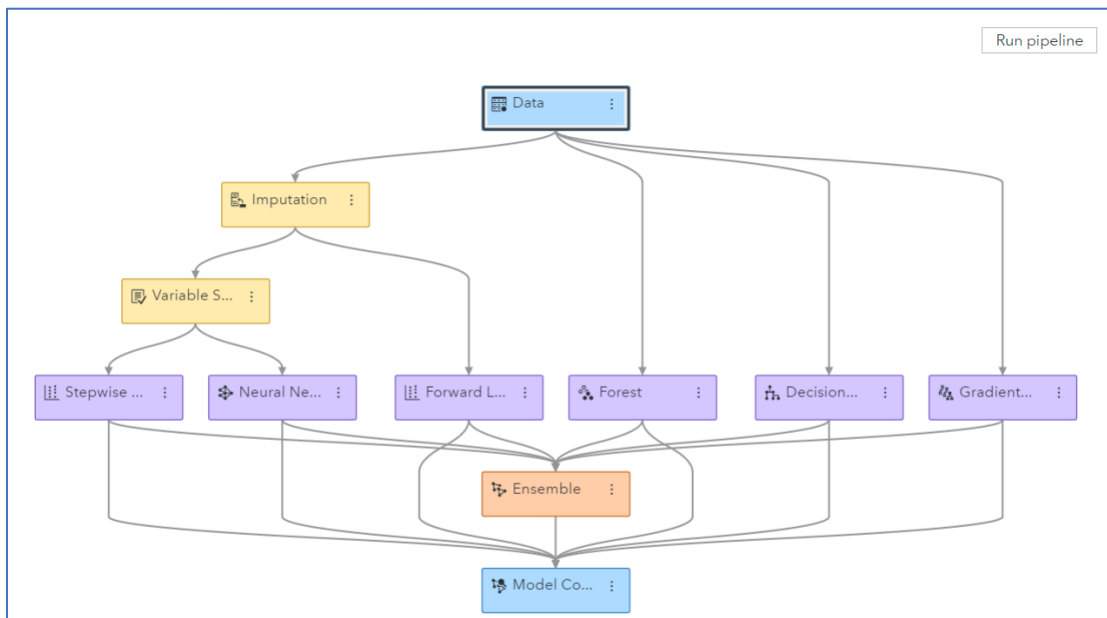
Template Name	Description	Owner	Last Modified
Advanced template for class target	Data mining pipeline that extends the intermediate template for a class target by adding neural network, forest, and gradient boosting models. An ensemble model is also provided.	SAS Pipeline	Sep 6, 2022, 4:25:44 PM
Advanced template for class target with autotuning	Data mining pipeline for a class target that contains autotuned tree, forest, neural network, and gradient boosting models.	SAS Pipeline	Sep 6, 2022, 4:25:42 PM
Advanced template for interval target	Data mining pipeline that extends the intermediate template for an interval target by adding neural network, forest, GAM, and gradient boosting models. An ensemble model is also provided.	SAS Pipeline	Sep 6, 2022, 4:25:51 PM
Advanced template for interval target with autotuning	Data mining pipeline that extends the intermediate template for an interval target by adding GAM and autotuned tree, forest, neural network, and gradient boosting models. An ensemble model is also provided.	SAS Pipeline	Sep 6, 2022, 4:25:47 PM
Basic template for class target	Data mining pipeline that contains a Data, Imputation, Logistic Regression, and Model Comparison node connected in a linear flow.	SAS Pipeline	Sep 6, 2022, 4:25:57 PM
Basic template for interval target	Data mining pipeline that contains a Data, Imputation, Linear Regression, and Model Comparison node connected in a linear flow.	SAS Pipeline	Sep 6, 2022, 4:26:05 PM

Cancel

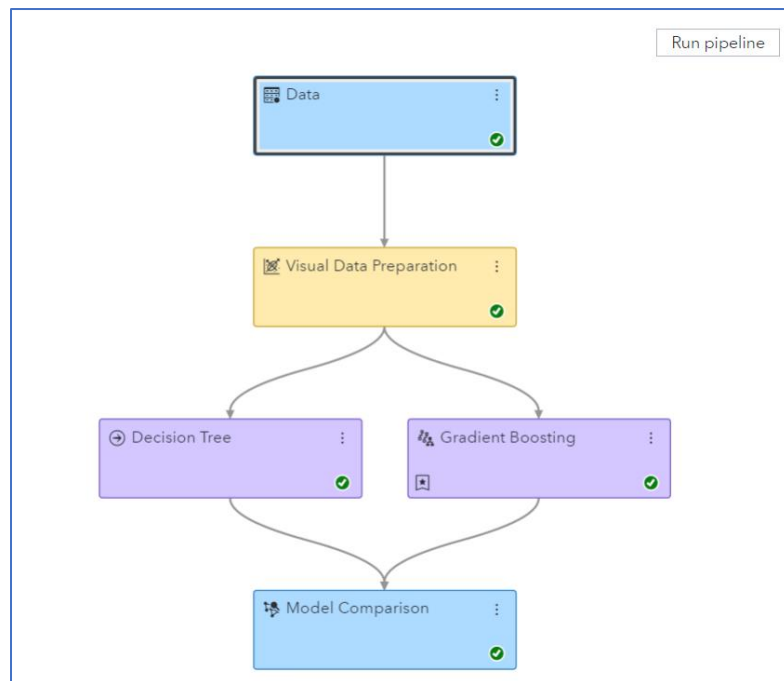
10. Since we have a class target (i.e., “Yes” or “No” for a bad loan), let’s get crazy and select **Advanced template for class target** and then click **OK**. The Advanced template contains a neural network, a stepwise logistic regression, a forward logistic regression, a decision tree, a gradient boosting model, and a forest model. Finally, it also contains an ensemble estimate of all those models put together.
11. We can then change the name of the pipeline to “Advanced Template” and then click **Save**:

The screenshot shows the 'New Pipeline' dialog box. The 'Name' field is highlighted with a red box and labeled '1', containing the text 'Advanced Template'. The 'Description' field is empty. The 'Select a pipeline template' section has a dropdown menu showing 'Advanced template for class target' and a 'Browse' button. The 'Automatically generate the pipeline' option is unchecked. The 'Set automation time limit' option is checked, with a value of '10' minutes. An 'Advanced Settings' button is visible. The 'Save' button is highlighted with a red box and labeled '2', and the 'Cancel' button is next to it.

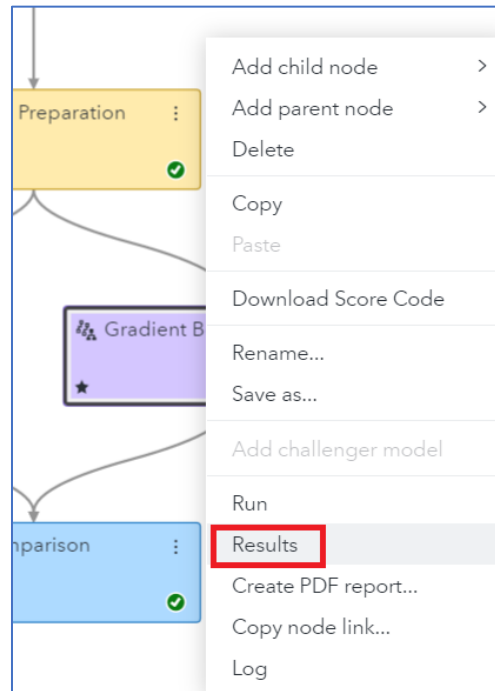
12. The new pipeline appears as follows:



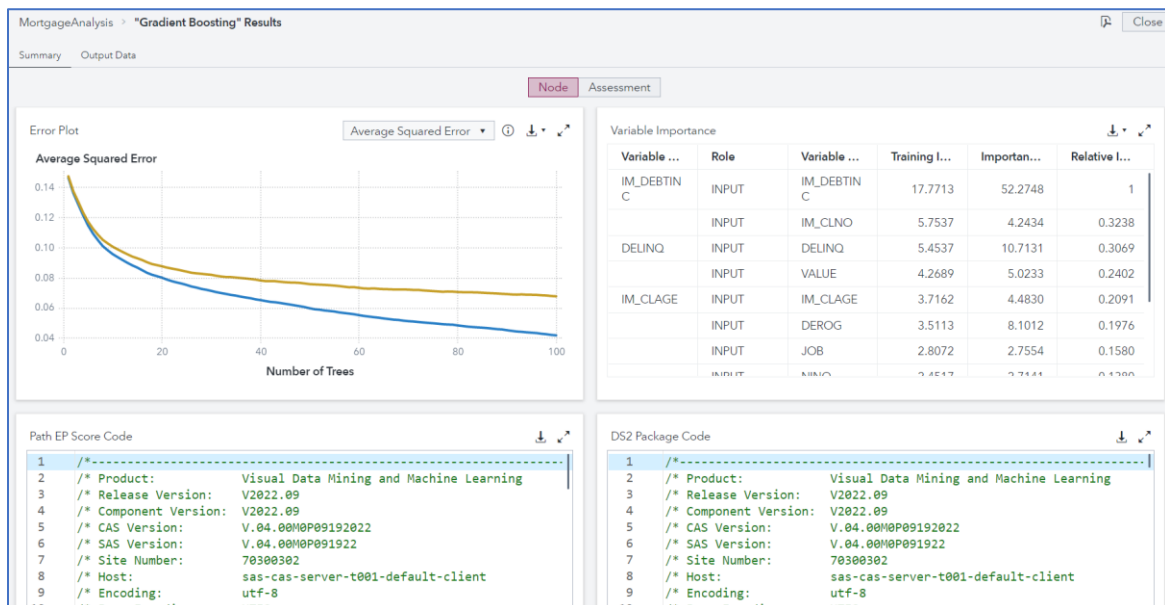
13. Again, look at all those modeling options! Yes, yes – that’s seven for the price of one (click)! Since we’re getting a bit pressed for time, let’s simply run all the models – by clicking **Run pipeline** – and return to our original pipeline, **Interactive-Model Pipeline**, to examine those results.
14. Green check marks in the pipeline indicate that the individual node has run. Thus, a completed **Interactive-Model Pipeline** appears as:



15. Let's examine our Gradient Boosting output. Right-click on the node, and select **Results**:



16. Explore both the **Node** and **Assessment** tabs so that you can better understand the results produced by SAS Model Studio:



17. Close **"GradientBoosting"Results** when you've had your fill.

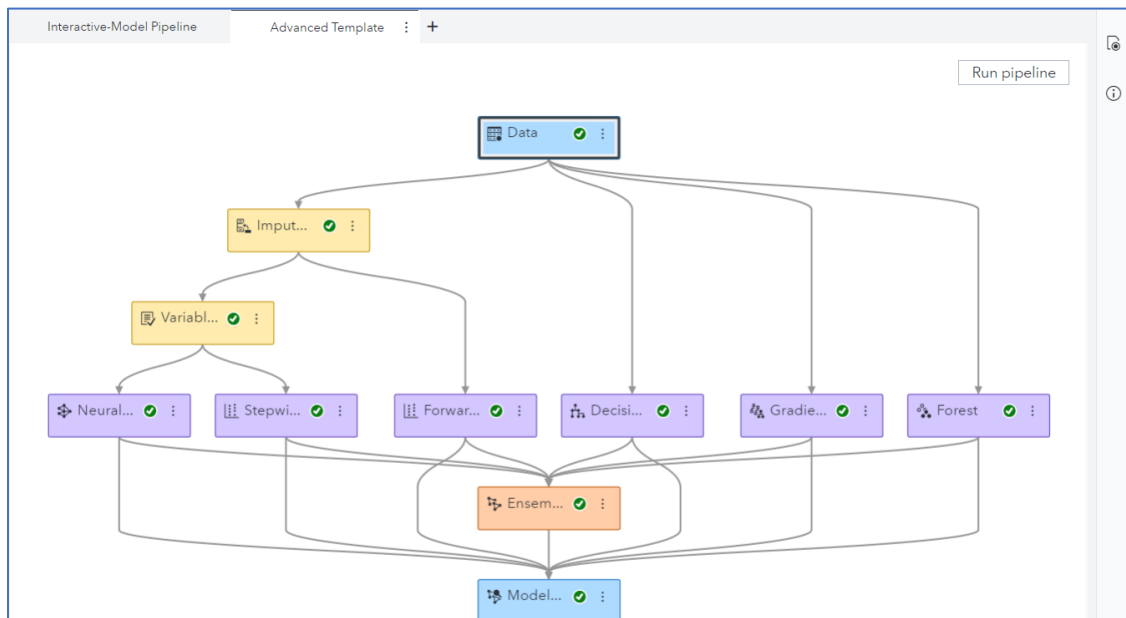
18. Next, select the Model Comparison node, right-click and select **Results**. Results will allow you to directly compare the two models in our pipeline:

Model Comparison													
Champi...	Name	Algorith...	KS (You...	Accuracy	Averag...	Area Un...	Cumula...	Cumula...	Cutoff	Data Role	Depth	F1 Score	False D
★	Gradient Boosting	Gradient Boosting	0.7342	0.9083	0.0678	0.9420	4.6779	46.7787	0.5000	VALIDATE	10	0.7429	0.15
	Decision Tree	Decision Tree	0.6228	0.8831	0.0925	0.8658	4.0816	40.8156	0.5000	VALIDATE	10	0.7027	0.28

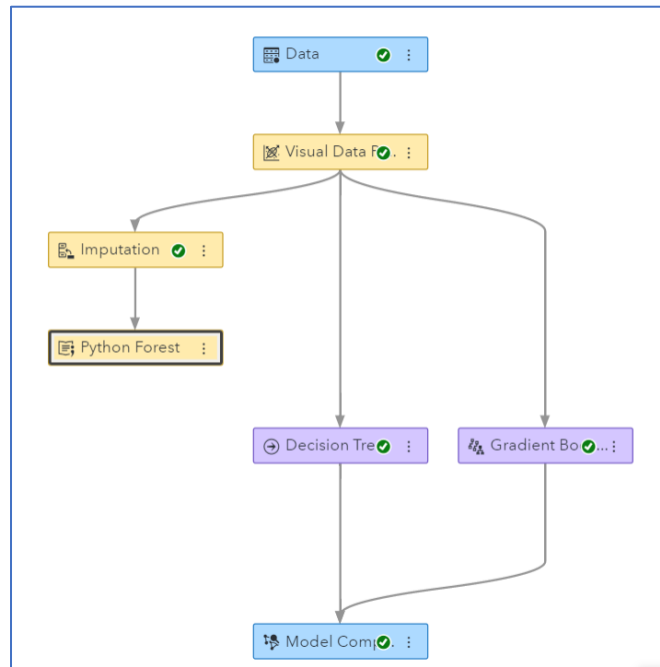
19. We can see that the Gradient Boosting model is selected as champion – because it has the highest KS(Youden) statistic. If you scroll to the right, you can see the Gradient Boosting model has a lower misclassification rate. Thus, the default Gradient Boosting model in SAS Model Studio – with all its hyper-tuned parameters – does a better job of modeling BAD=1 than the decision tree used in both SAS VA and SAS Model Studio *and* the default gradient boosting model used in SAS VA. So, it looks like moving the modeling to SAS Model Studio improved our results. Good stuff!

20. After processing all the output from our first pipeline – we’ve nearly forgotten: we have a second pipeline to examine! Let’s **Close** the Model Comparison results window and return to the **Advanced Template**.

21. Ensure that the Advanced Template is finished running – which is indicated by a bunch of green checkmarks:



22. Check, check, check. Let’s go straight to the Model Comparison and find which model is champion. Again, right-click on the **Model Comparison** node, select **Results**, and then explore the following:



29. Under **Node options** for our **Python Forest**, verify that the language is set to **Python**. It should be Python by default...
30. Expand the **Data Sample** properties. Clear the check box for **Include SAS formats**.
 - a. A fun note: this property controls whether the downloaded data sent to Python or R should keep SAS formats. Why? Well, it is usually recommended that you keep SAS formats, and this should work in most cases. But... some numeric formats such as DOLLARw.d add a dollar sign and change the data type of the variable when exporting to CSV. In such cases, these formats must be removed.
31. Also change the **Sampling method:** to **(none)**. The cumulative changes:

>>

Python Forest 🔍 ▶ 🖨️ ?

Description:

Runs Python or R code.

Open code editor

Language:

Python ▼

▼ Input to Open Source

▼ Data Sample

Sampling method:

(none) ▼

☐ Include SAS formats

☒ Generate data frame

32. Now, click the **Open code editor** button in the Node Options pane:

Open Code Editor

33. Now it's time to copy-and-paste some Python Code. Grab the following:

```
from sklearn import ensemble

# Get full data with inputs + partition indicator
dm_input.insert(0, dm_partitionvar)
fullX = dm_inputdf.loc[:, dm_input]

# Dummy encode class variables
fullX_enc = pd.get_dummies(fullX, columns=dm_class_input, drop_first=True)

# Create X (features/inputs); drop partition indicator
X_enc = fullX_enc[fullX_enc[dm_partitionvar] == dm_partition_train_val]
X_enc = X_enc.drop(dm_partitionvar, 1)

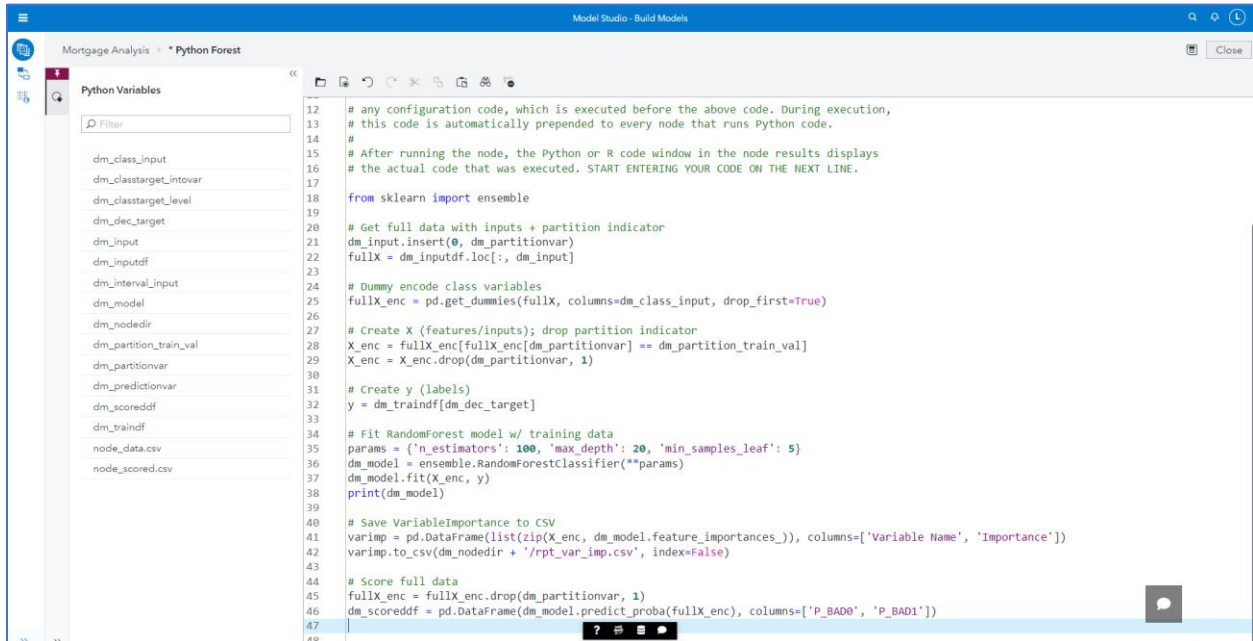
# Create y (labels)
y = dm_traindf[dm_dec_target]

# Fit RandomForest model w/ training data
params = {'n_estimators': 100, 'max_depth': 20, 'min_samples_leaf': 5}
dm_model = ensemble.RandomForestClassifier(**params)
dm_model.fit(X_enc, y)
print(dm_model)

# Save VariableImportance to CSV
varimp = pd.DataFrame(list(zip(X_enc, dm_model.feature_importances_)),
columns=['Variable Name', 'Importance'])
varimp.to_csv(dm_nodedir + '/rpt_var_imp.csv', index=False)

# Score full data
fullX_enc = fullX_enc.drop(dm_partitionvar, 1)
dm_scoreddf = pd.DataFrame(dm_model.predict_proba(fullX_enc),
columns=['P_BAD0', 'P_BAD1'])
```

34. Yes – that’s a lot of code. Now paste it into the Python Forest window:



```

12 # any configuration code, which is executed before the above code. During execution,
13 # this code is automatically prepended to every node that runs Python code.
14
15 # After running the node, the Python or R code window in the node results displays
16 # the actual code that was executed. START ENTERING YOUR CODE ON THE NEXT LINE.
17
18 from sklearn import ensemble
19
20 # Get full data with inputs + partition indicator
21 dm_input.insert(0, dm_partitionvar)
22 fullX = dm_inputdf.loc[:, dm_input]
23
24 # Dummy encode class variables
25 fullX_enc = pd.get_dummies(fullX, columns=dm_class_input, drop_first=True)
26
27 # Create X (features/inputs); drop partition indicator
28 X_enc = fullX_enc[fullX_enc[dm_partitionvar] == dm_partition_train_val]
29 X_enc = X_enc.drop(dm_partitionvar, 1)
30
31 # Create y (labels)
32 y = dm_traindf[dm_dec_target]
33
34 # Fit RandomForest model w/ training data
35 params = {'n_estimators': 100, 'max_depth': 20, 'min_samples_leaf': 5}
36 dm_model = ensemble.RandomForestClassifier(**params)
37 dm_model.fit(X_enc, y)
38 print(dm_model)
39
40 # Save VariableImportance to CSV
41 varimp = pd.DataFrame(list(zip(X_enc, dm_model.feature_importances_)), columns=['Variable Name', 'Importance'])
42 varimp.to_csv(dm_nodedir + '/rpt_var_imp.csv', index=False)
43
44 # Score full data
45 fullX_enc = fullX_enc.drop(dm_partitionvar, 1)
46 dm_scoreddf = pd.DataFrame(dm_model.predict_proba(fullX_enc), columns=['P_BAD0', 'P_BAD1'])
47
48

```

35. A few notes about the Python code – while saving more of the gory details for another time:

- This code fits a random forest classifier model in Python. The default values for the parameters that control the size of the trees (for example, **max_depth (default=None)**, **min_samples_leaf (default=1)**) lead to fully grown and unpruned trees, which can be very large data sets. To reduce memory consumption, the complexity and size of the trees are controlled by setting parameter values like the ones in the code above.
- The code that needs to be changed for different data sets is the last line – i.e., naming your predictions with the **P_ + “target”** naming convention.
- It’s important to note that we are just modeling the data here. Currently, there is not a way to do data preparation within the Open Source Code node so that a subsequent node will recognize it. If this is necessary, either prepare data before Model Studio, or perform both of the following: (1) open-source data preparation with the Open Source Code node (in the Preprocessing group), and (2) modeling with the Open Source Code node (in the Supervised Learning group).

36. In the upper right corner of the window, click the **Save** icon to save the Python code and then click the **Close** button to close the Code Editor window.

37. There is one more setting to change. Select the **Use output data in child nodes** property under **Environment**:

Python Forest

Description:

Runs Python or R code.

Open code editor

Language:

Python

Input to Open Source

Data Sample

☒ Generate data frame

Environment

☒ Use output data in child nodes

38. A little bit behind the why of the *Use output data in child nodes* button. Every time that this property is set, a copy of the output data is saved in the Model Studio project library, which can be used in later stages of the pipeline.

39. Now it is time to run the **Python Forest** node. Right-click on the node and select **Run**.

40. Open the **Results** of the Python Forest node to examine output.

Mortgage Analysis > "Python Forest" Results

Summary

Output Data

rpt_var_imp.csv

Variable Name	Training Importance
IMP_VALUE	0.0703
IM_CLAGE	0.0873
IM_DEBTINC	0.2938
IM_MORTDUE	0.0533
LOAN	0.0848
IMP_DELIQ_1	0.0277
IMP_DELIQ_2	0.0242
IMP_DELIQ_3	0.0322

Python Code

```

1 #-----
2 # Language: PYTHON
3 #-----
4 # Variable declarations
5 dm_nodedir = '/opt/sas/v4e055/config/var/tmp/compsrv/default/7fec17fb-ef2c-4e4c-9b0d-0f511dc
6 dm_dec_target = 'BAD'
7 dm_partitionvar = 'PartInd'
8 dm_partition_train_val = 1
9
10 dm_class_input = ['IMP_DELIQ', 'IMP_DEROG', 'IMP_JOB', 'IMP_NINQ', 'IMP_REASON', 'IM_CLNO', 'IM_Y
11 dm_interval_input = ['IMP_VALUE', 'IM_CLAGE', 'IM_DEBTINC', 'IM_MORTDUE', 'LOAN']
12 dm_input = dm_class_input + dm_interval_input
13 dm_predictionvar = ['P_BAD0', 'P_BAD1']
14 dm_classtarget_intovar = 'I_BAD'
15

```

Python Output

```

1 RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
2 criterion='gini', max_depth=20, max_features='auto',
3 max_leaf_nodes=None, max_samples=None,
4 min_impurity_decrease=0.0, min_impurity_split=None,
5 min_samples_leaf=5, min_samples_split=2,
6 min_weight_fraction_leaf=0.0, n_estimators=100,
7 n_jobs=None, oob_score=False, random_state=None,
8 verbose=0, warm_start=False)
9

```

Data Preparation Code

```

1 package MS_1de1991fcd1a449c9dd30233e4345460_18MAY2023131834314 / overwrite=yes;
2 dcl double "IMP_DELIQ" having label n'Imputed DELIQ';
3 dcl double "IMP_DEROG" having label n'Imputed DEROG';
4 dcl nchar(7) "IMP_JOB" having label n'Imputed JOB';
5 dcl double "IMP_NINQ" having label n'Imputed NINQ';
6 dcl nchar(7) "IMP_REASON" having label n'Imputed REASON';
7 dcl double "IMP_VALUE" having label n'Imputed VALUE';
8 varlist allvars [_all_];
9
10 method _EHG0MS3WC036Y0MS063C3990P();
11
12 if MISSING("DELIQ") then do ;
13 "IMP_DELIQ" = 0.0;
14 end;
15

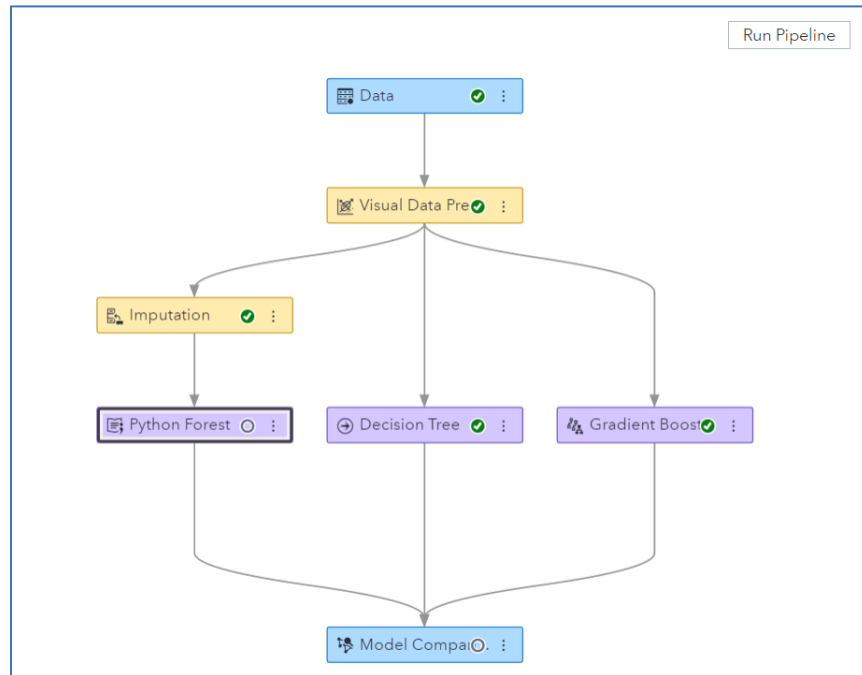
```

41. Examine the results for a bit, so you get a better understanding of model performance. Click **Close** when you're satiated.

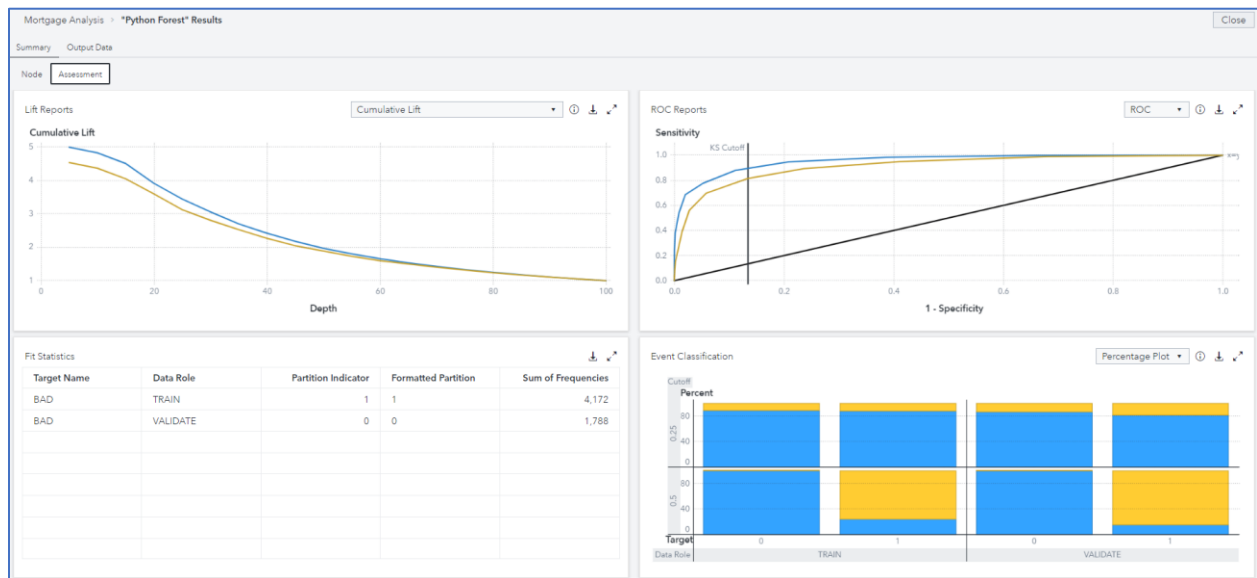
42. And after you're done exploring, I have a thought-provoking question for you: *Why does the Open Source Code node lack assessment results even though it ran successfully?*

Well, for model assessment results, you need to move the nodes to the supervised learning lane.

43. To get our open-source node registered as a model, right-click the **Python Forest** node and select **Move** ⇒ **Supervised Learning**. Your new view should appear as follows:



44. The color of *Python Forest* has changed to purple, showing that the node is now part of the Supervised Learning group. Notice also that the nodes need to be rerun, as the green check mark has disappeared.
45. Click the **Run Pipeline** button.
46. Open the **Results** of the Python Forest. Click the **Assessment** tab to learn even more about model performance:



47. The usual assessment results are displayed. Explore until your heart is content and then close the results.


48. Open the **Results** of the **Model Comparison** node. Which model is our ultimate champion in the *Interactive-Model Pipeline*?

Champion	Name	Algorithm Name	KS (Youden)	Misclassification Rate
★	Gradient Boosting	Gradient Boosting	0.7195	0.0906
	Python Forest	Open Source Code	0.6734	0.1683
	Decision Tree	Decision Tree	0.6319	0.1331


49. For our modeling, Gradient Boosting algorithm has the highest KS (Youden) statistic, followed closely by the Python Forest. So, it looks like both additions improve the fit over our default Decision Tree from SAS VA. Woot!

50. **Close** out of the Model Comparison Results window.

51. Let's do two more things before we wrap up. The first is to examine the Pipeline Comparison tab, found here:

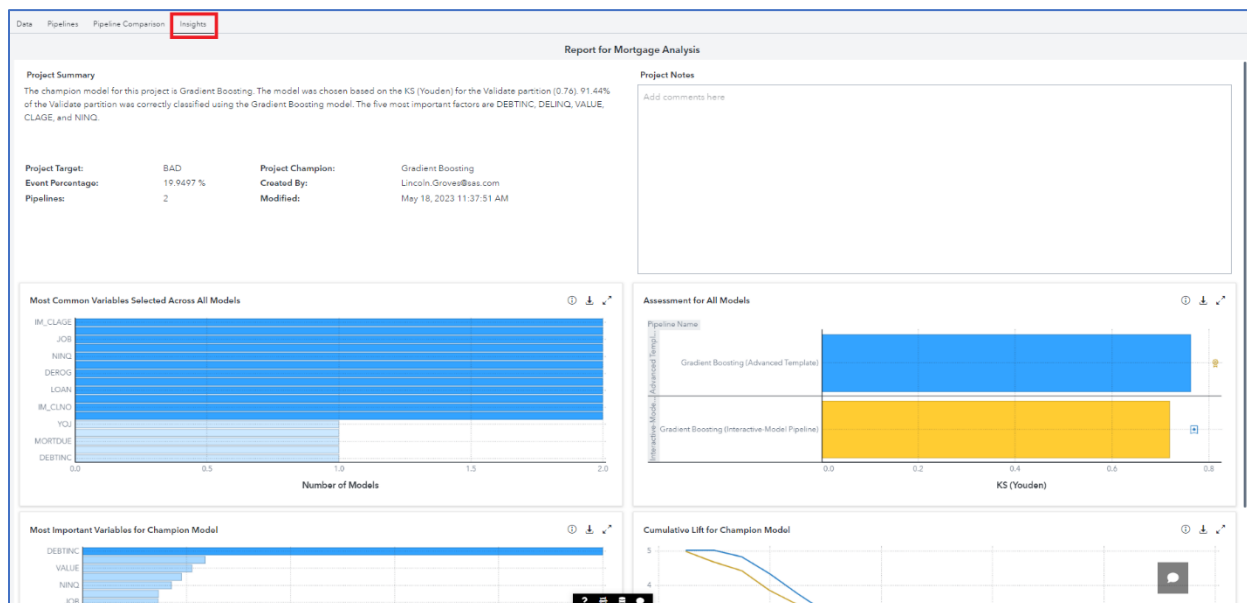


52. Rather than examining the champion of an individual pipeline, this will give us a champion of all the pipelines. We have only two potential models to choose from – but we could have many, many more! Select **Pipeline Comparison** tab – and examine the following:

<input type="checkbox"/>	Champion	Name	Algorithm Name	Pipeline Name	KS (Youden)
<input checked="" type="checkbox"/>		Gradient Boosting	Gradient Boosting	Advanced Template	0.763
<input type="checkbox"/>		Gradient Boosting	Gradient Boosting	Interactive-Model Pipeline	0.719

53. No surprises here – a Gradient Boosting model is the overall champion. But it’s the one configured in the **Advanced Template** that wins the crown, with a slightly higher KS (Youden) statistic. *And, as stated earlier, your results may differ marginally due to different random samples being used.*

54. Finally, let’s click on the **Insights** tab, found to the right of Pipeline Comparison:



55. Storytelling is an important part of the data analytics process, as you communicate your findings with others. The **Insights** tab contains several helpful resources – from overall project summaries to assessments of models and variables – that can help you better explain your analytic process and overall champion model.

56. Explore this output at your leisure. Then smile – you’ve made it through your first day at iLink Mortgage, Inc!

Appendix

Appendix A: Access Software

Steps to accessing SAS Viya for Learners for the first time:

1. Navigate to the SAS Viya for Learners webpage,
https://www.sas.com/en_us/software/viya-for-learners.html
2. Click the "Access for Educators" or "Access for Students" button based on your role.
3. Log in with your SAS Profile that is linked to an academic affiliation. If you don't have a SAS Profile, click [here](#) to set one up.
4. Access the software by clicking on the **Launch SAS Viya for Learners** button.


For future sessions with VFL, you can access VFL by visiting <https://vle.sas.com/vfl>.

Contact us at academic@sas.com if you have any questions about access.

Appendix B: Helpful Documentation

Here are two options for additional help and guidance:

1. **SAS Help Center** in SAS VFL
 - a. It's been lurking in the upper-right-hand corner all long!



- b. Click the button to go to the **SAS Help Center** and the **SAS Documentation**.
2. **SAS Video Library**
 - a. Prefer videos instead? We've got you covered here: <https://video.sas.com/>
 - b. Check out the **How-to Tutorials** or simply **Search Videos**.

Appendix C: Recommended Learning

The [SAS Global Academic Program](#) offers free eLearning courses for students to learn SAS through the Student Skill Builder. The following eLearning courses and paths available are recommended to help with this activity:

- SAS Visual Analytics 1 for SAS Viya: Basics
- SAS Visual Analytics 2 for SAS Viya: Advanced
- SAS Visual Statistics in SAS Viya: Interactive Model Building
- Machine Learning Using SAS Viya
- Programming for SAS Viya

Alternatively, the [SAS Learning Subscription](#) grants you access to an extensive library of SAS eLearning courses. Sign up for a free 7-day trial [here](#)!